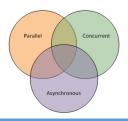


Agenda



What is concurrency - parallelism - asynchronous - multithreading

Different programming models (pro & cons) – with code sample

Code Examples

Threading & Parallel Programming

Asynchronous Workflows

Actor model

Alternative STM - Channel ...

Introduction - Riccardo Terrell

- □ Originally from Italy, currently Living/working in Washington DC ~10 years
- \Box +/- 19 years in professional programming
 - C++/VB \rightarrow Java \rightarrow .Net C# \rightarrow Scala \rightarrow Haskell \rightarrow C# & F# \rightarrow ??
- DC F# User Group Organizer
- Working @ III statmuse
- Polyglot programmer believes in the art of finding the right tool for the job







Moore's law - The Concurrency challenge



Modern computers come with multiple processors, each equipped with multiple cores, but the single processor is slower than it used to be!

Windows Task No le Options View	-		13		_(5)
Applications Process	ses Services Pe	formance Net	working Users		
CPU Usage	CPU Usage His	tory	-		
99 %					
Memory	Physical Memor	y Usage History		 	
22.1 68					









Multithreading in practice



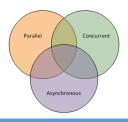






- Shared Memory Concurrency
- Data Race / Race Condition
- Works in sequential single threaded environment
- Very difficult to parallelize
- Difficult to maintain and test
- Locking, blocking, call-back hell

Goals



Embrace Parallelism
exploiting the
potential of multicore
CPUs

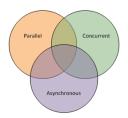
No single solution...

You must combine different programming models

(right tool for the job)

FP really **shines** in the area of concurrency



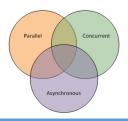


Concurrent

Multithreaded

Parallel

Asynchronous



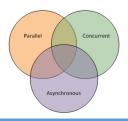
Concurrent

Multithreaded

Parallel

Asynchronous

Several things happening at once TPL — Async — Agent - RX



Concurrent

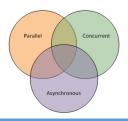
Multithreaded

Parallel

Asynchronous

Several things happening at once TPL – Async – Agent - RX

Multiple execution context



Concurrent

Multithreaded

Parallel

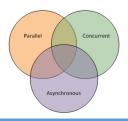
Asynchronous

Several things happening at once TPL — Async — Agent - RX

Multiple execution context

Multiple simultaneous computations

TPL - Async - Agent - RX - GPU



Concurrent

Multithreaded

Parallel

Asynchronous

Several things happening at once TPL – Async – Agent - RX

Multiple execution context

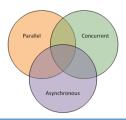
Multiple simultaneous computations

TPL - Async - Agent - RX - GPU

Asynchronous
AsyncWorkflow – TAP – Agent

Concurrent Programming Models

Sequential Fuzzy Match

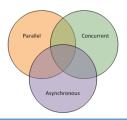


```
List<string> matches = new List<string>();
foreach (var word in WordsToSearch)
{
    var localMathes = JaroWinklerModule.bestMatch(Words, word);
    matches.AddRange(localMathes.Select(m => m.Word));
}
```

Fuzzy match 7 words against 13.4 Mb of text

Time execution in 4 Logical cores – 6 Gb Ram: **23,167** ms

Two Threads Fuzzy Match



```
List<string> matches = new List<string>();
    var t1 = new Thread(() =>{
        var take = WordsToSearch.Count / 2;
        foreach (var word in WordsToSearch.Take(take))
            var localMathes = FuzzyMatch.JaroWinklerModule.bestMatch(Words, word);
            matches.AddRange(localMathes.Select(m => m.Word));
    });
    var t2 = new Thread(() =>{
        var start = WordsToSearch.Count / 2;
        var take = WordsToSearch.Count - start;
        foreach (var word in WordsToSearch.Skip(start).Take(take)) {
            var localMathes = FuzzyMatch.JaroWinklerModule.bestMatch(Words, word);
            matches.AddRange(localMathes.Select(m => m.Word));
    });
    t1.Start();
                                     t2.Start();
                                     t2.Join();
    t1.Join();
```

Time execution: **15,436** ms



Multi Thread Fuzzy Motch

```
List<string> matches = new List<s
var threads = new Thread[Environn
                                WRONGIA
for (int i = 0; i < threads.Lengt</pre>
      threads[index,
      var take = Word
                                                                            wads.Length));
      var start =
                                                                                index * take;
                                                               e)) {
      foreach (var
                 var log
                                                  bWink
                                                                e.bestMatch(Words, word);
                 matc
                                                  ect(m = 1)
                                                               ((rd)
 for (int i = 0: i < threads.Length; i++)</pre>
    threads[i].Start();
for (int i = 0: i < threads.Length; i++)</pre>
   threads[i].Join();
```

Time execution: **6,857** ms



Multi Thread Fuzzy Match

```
List<string> matches = new List<string>():
var threads = new Thread[Environment.ProcessorCount];
for (int i = 0; i < threads.Length; i++) {</pre>
                  var index = i:
     threads[index] = new Thread(() => {
     var take = WordsToSearch.Count / (Math.Min(WordsToSearch.Count. threads.Length));
     var start = index == threads.Length - 1 ? WordsToSearch.Count - take : index * take;
     foreach (var word in WordsToSearch.Skip(start).Take(take)) {
               lock (matches)
                     macches.AudRange(localMathes.Select(m => m.Word));
                  });
for (int i = 0; i < threads.Length; i++)</pre>
   threads[i].Start();
for (int i = 0; i < threads.Length; i++)</pre>
  threads[i].Join();
```

Time execution: **7,731** ms

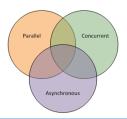
Task Parallel Library (TPL) - Task Programming





- Efficient and scalable use of system resources...
 - When used correctly
- Programmatic fine grain control
- Model largely applicable
- Easy to integrate into existing programs
- Hard to solve complex problems
- Sometimes requires the introduction of locking primitives

Parallel Loop Fuzzy Match

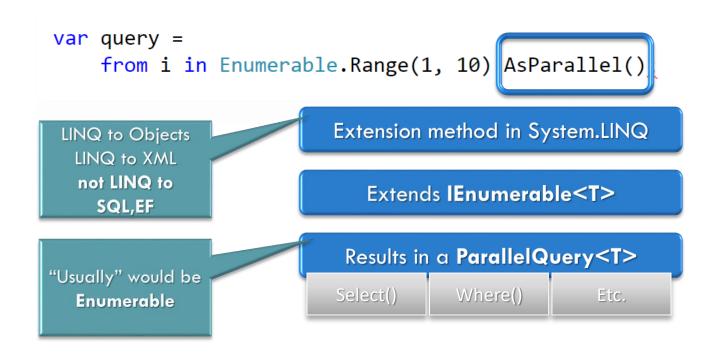


```
List<string> matches = new List<string>();
object sync = new object();
Parallel.ForEach(WordsToSearch,
                 // thread local initializer
                  () => { return new List<string>(); },
                  (word, loopState, localMatches) => {
                         var localMathes = FuzzyMatch.JaroWinklerModule.bestMatch(Words, word);
                         localMatches.AddRange(localMathes.Select(m => m.Word));// same code
                         return localMatches;
     (finalResult) =>
        // thread local aggregator
        lock (sync) matches.AddRange(finalResult);
```

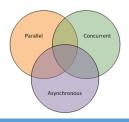
Time execution: **7,429** ms

Parallel LINQ (PLINQ)



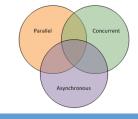


PLINQ Fuzzy Match

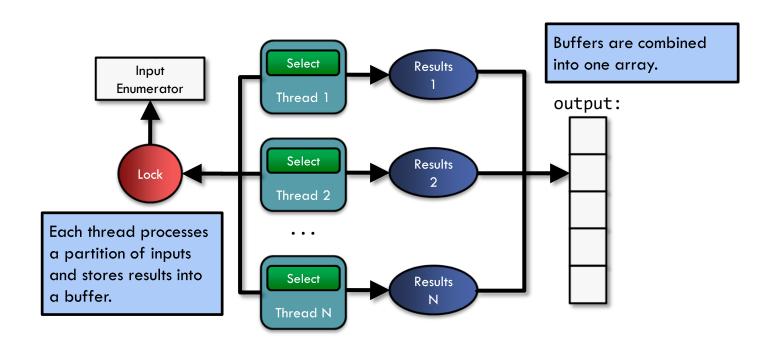


```
ParallelQuery<string> matches =
    (from word in WordsToSearch.AsParallel()
    from match in JaroWinklerModule.bestMatch(Words, word)
    select match.Word);
```

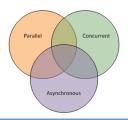
Time execution: **6,347** ms



PLINQ under the hood

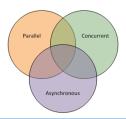


PLINQ options



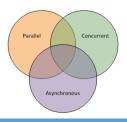
```
var parallelQuery = from t in source.AsParallel()
    select t:
var cts = new CancellationTokenSource();
cts.CancelAfter(TimeSpan.FromSeconds(3));
    parallelQuery
        .WithDegreeOfParallelism(Environment.ProcessorCount)
        .WithExecutionMode(ParallelExecutionMode.ForceParallelism)
        .WithMergeOptions(ParallelMergeOptions.Default)
        .WithCancellation(cts.Token)
        .ForAll(Console.WriteLine);
```

Multi Task Fuzzy Match



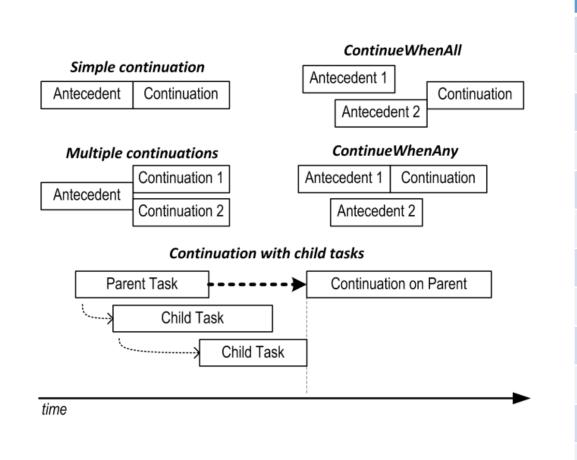
```
var tasks = new List<Task<List<string>>>();
var matches = new ThreadLocal<List<string>>(() => new List<string>());
foreach (var word in WordsToSearch) {
        tasks.Add(Task.Factory.StartNew<List<string>>((w) => {
         List<string> localMatches = matches.Value;
         var localMathes = FuzzyMatch.JaroWinklerModule.bestMatch(Words, w);
        localMatches.AddRange(localMathes.Select(m => m.Word));
         return localMatches;
        }, word));
Task.Factory.ContinueWhenAll(tasks.ToArray(), (ts) =>
      return new List<string>(tasks.SelectMany(t => t.Result).Distinct())
).Wait();
                                                   Time execution: 4,192 ms
```

PLINQ Fuzzy Match



Create a load-balancing partitioner, specify false for static partitioning.

Time execution: **4,217** ms



TaskContinuationOptions

AttachedToParent

ExecuteSynchronously

LazyCancellation

LongRunning

None

NotOnCanceled

NotOnFaulted

NotOnRanToCompletion

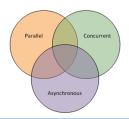
OnlyOnCanceled

OnlyOnFaulted

OnlyOnRanToCompletion

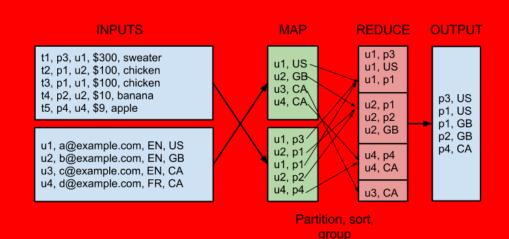
Run Continuations A synchronously

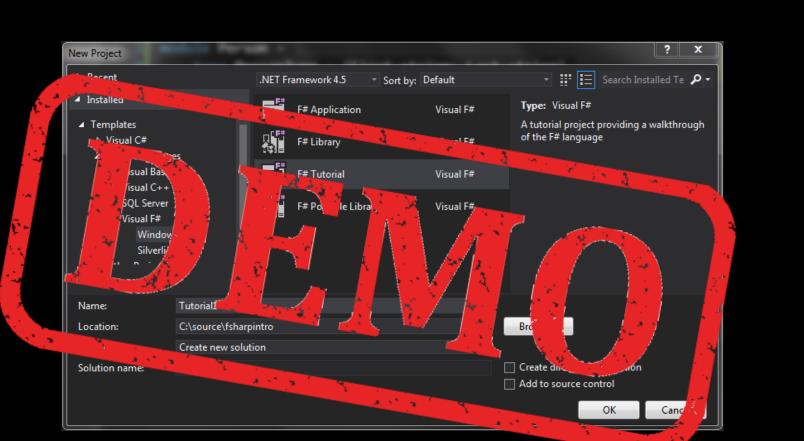
Task continuation



```
task1.ContinueWith(antecedent =>
        Console.WriteLine("Task #1 completion continuation."),
        TaskContinuationOptions.OnlyOnRanToCompletion);
task1.ContinueWith(antecedent =>
        Console.WriteLine("Task #1 cancellation continuation."),
        TaskContinuationOptions.OnlyOnCanceled);
task1.ContinueWith(antecedent =>
        Console.WriteLine("Task #1 on error continuation."),
        TaskContinuationOptions.OnlyOnFaulted);
task1.ContinueWith(antecedent =>
        Console.WriteLine("Task #1 continuation lomg running."),
        TaskContinuationOptions.LongRunning);
```

Data Parallelism & Lambda Architecture

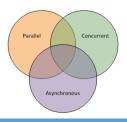




Functional Programming in Concurrency

- declarative
- functions as values
- side-effects free

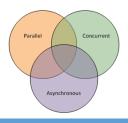
- referential transparency
- immutable
- composition



Immutability

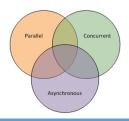
```
public class ImmutablePerson
    public ImmutablePerson(string firstName, string lastName, int age) {
        FirstName = firstName;
        LastName = lastName;
        Age = age;
    public string LastName { get; }
    public string FirstName { get; }
    public int Age { get; }
    public ImmutablePerson UpdateAge(int age)
        return new ImmutablePerson(this.FirstName, this.LastName, age);
```

type Person = {FirstName:string; LastName:string; Age:int}



Compare And Swap - CAS

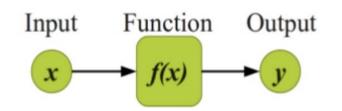
```
void DoWithCAS<T>(ref T location, Func<T, T> generator) where T : class
{
    T temp, replace;
    do
    {
        temp = location;
        replace = generator(temp);
    } while (Interlocked.CompareExchange(ref location, replace, temp) != temp);
}
```

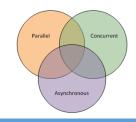


Compare And Swap - CAS

```
public class LockFreeStack<T>
   private class Node {
       public T Data;
        public Node Next;
   private Node head;
   public void Push(T element)
        Node node = new Node {Data = element};
       DoWithCAS(ref head, h =>
            node.Next = h;
            return node;
        });
```

Side Effects





```
int WordCounterLing(string dirPath) {
    var wordTable= (from file in Directory.EnumerateFiles(dirPath, "*.*")
                       from line in File.ReadAllLines(file)
                       from word in Regex.Split(line, @"\W+")
                       where !string.IsNullOrEmpty(word) && word.Length > 0
                       select word.ToUpper())
                  .GroupBy(s => s).ToDictionary(k => k.Key, v => v.Count());
    return wordTable;
```

Input Function Output f(x)Asynchro

Concurrent

Side Effects Free

Asynchronous Programming

Asynchronous Programming



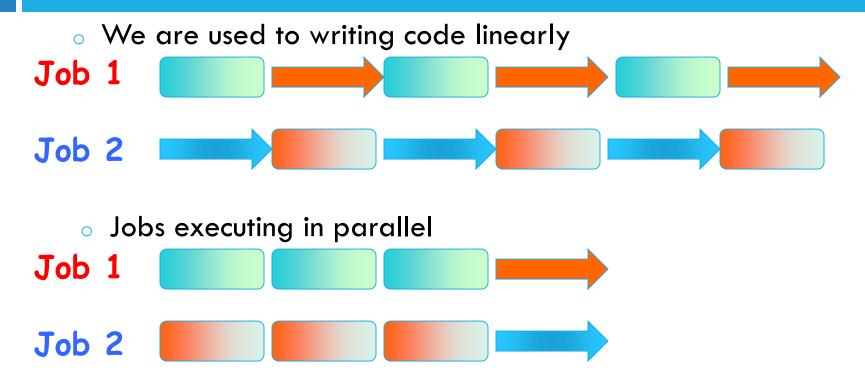
- Software is often I/O-bound, it provides notable performance benefits
 - Connecting to the Database
 - Leveraging web services
 - Working with data on disk
- Network and disk speeds increasingly slower
- □ Not Easy to predict when the operation will complete (no-deterministic)
- IO bound functions can scale regardless of threads
 - IO bound computations can often "overlap"
 - This can even work for huge numbers of computations



Classic Synchronous programming

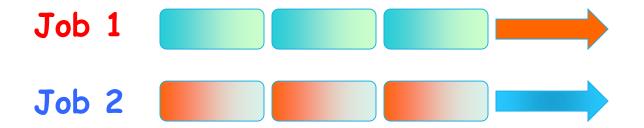
We are used to writing code linearly
Job 1
Job 2

Classic Synchronous programming



Classic Asynchronous programming





- Classic Asynchronous programming requires decoupling Begin from End
- Very difficult to:
 - Combine multiple asynchronous operations
 - Deal with exceptions, cancellation and transaction

Classic Asynchronous Programming



```
let callBack (callBack: JAsyncRes
  let fs = callBack.As
                 Where is my DATA?

Where is my DATA?
  fs.EndRead(callBa
  fc Dichago
let fs
  let asyncP
            ...ead(d
                                   ength,
    fs.Be
```

.0

Asynchronous Programming in C# 5.0

```
Task<string> GetTitleTplAsync(string url) {
           var w = new WebClient();
           Task<string> contentTask = w.DownloadStringTaskAsync(url);
           return contentTask.ContinueWith(t => {
               string result = ExtractTitle(t.Result);
               w.Dispose();
               return result; });
async Task<string> GetTitleCsAsync(string url) {
             using (var w = new WebClient()){
                 string content = await w.DownloadStringTaskAsync(url);
                 return ExtractTitle(content);
```

Anatomy of Async Workflows F#



```
let readData path = async {
  let stream = File.OpenRead(path)
  let! data = stream.AsyncRead(stream.Length)
  return data }
```

- Easy transition from synchronous
 - Wrap in asynchronous workflow with the *async* keyword, use *let!* for async calls and add *return*
 - No need for explicit callback
 - Easy to debug
- □ Supports loops, recursion, exceptions, cancellation, resource management
- Operation complete in the same scope

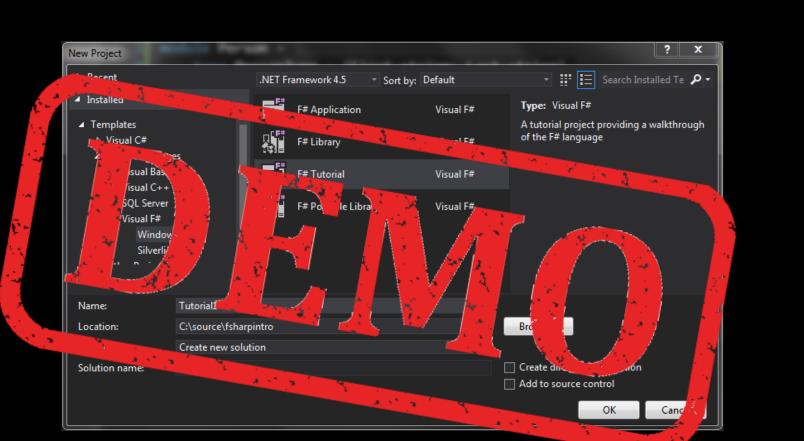
Anatomy of Async Workflows



```
let readData path : Async<byte[]> = async {
  let stream = File.OpenRead(path)
  let! data = stream.AsyncRead(stream.Length)
  return data }
```

Async defines a block of code which executes on demand

Easy to compose

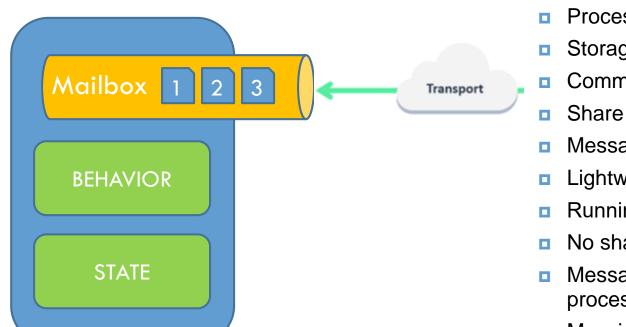


Message Passing

Agent (and Actor) model

Message Passing based concurrency

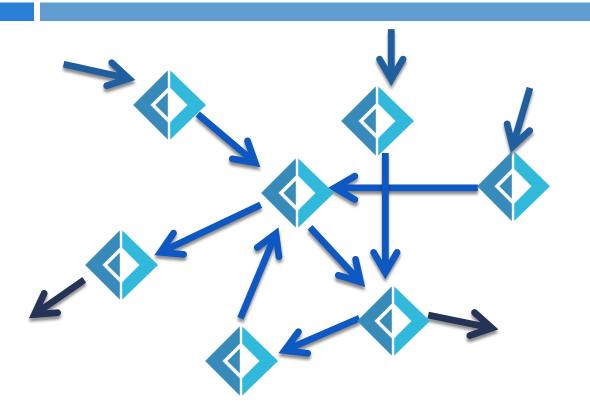




- Processing
- Storage State
- Communication only by messages
- Share Nothing
- Message are passed by value
- Lightweight object
- Running on it's own thread
- No shared state
- Messages are kept in mailbox and processed in order
- Massively scalable and lightening fast because of the small call stack

Message Passing based concurrency





Agents exchange messages

Receive message and react

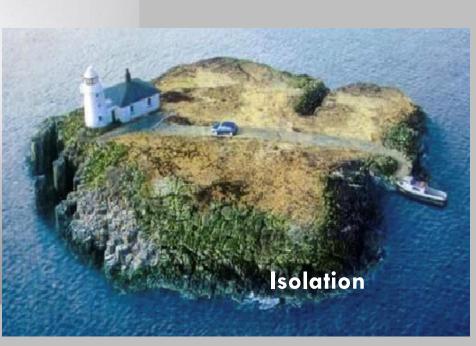
Reactive system

Handle inputs while running Emits results while running

The Solution is Immutability and Isolation

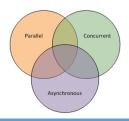






Immutability OR Isolation



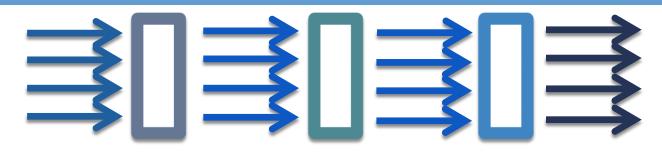


TPL DataFlow

```
var InputBlock = new BufferBlock<Tuple<string, string>>():
var splitLines = new TransformBlock<Tuple<string, st</pre>
        n => /* Code ... Transfirmation */ );
var splitWords = new TransformBlock<Tuple<string, st</pre>
        n => /* Code ... Transfirmation */ );
var fuzzyMatch = new TransformBlock<Tuple<string, st</pre>
        n => /* Code ... Transfirmation */ );
var finalBlocl = new ActionBlock<Tuple<string, strin</pre>
        n = /* Code ... */ );
InputBlock.LinkTo(splitLines, new DataflowLinkOption
splitLines.LinkTo(splitWords, new DataflowLinkOption
InputBlock.SendAsync(message);
```

Platform	Dataflow support
.NET 4.5	✓
.NET 4.0	×
Mono iOS/Droid	✓
Windows Store	~
Windows Phone Apps 8.1	✓
Windows Phone SL 8.0	~
Windows Phone SL 7.1	×
Silverlight 5	×

Pipeline Processing

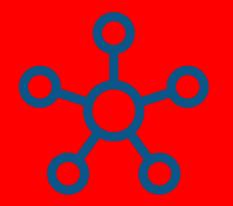


- Pipeline according to Wikipedia:
 - A pipeline is a set of data processing elements connected in series, so that the output of one element is the input of the next one. The elements of a pipeline are often executed in parallel or in time-sliced fashion; in that case, some amount of buffer storage is often inserted between elements

Agent Stock Ticker

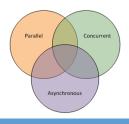


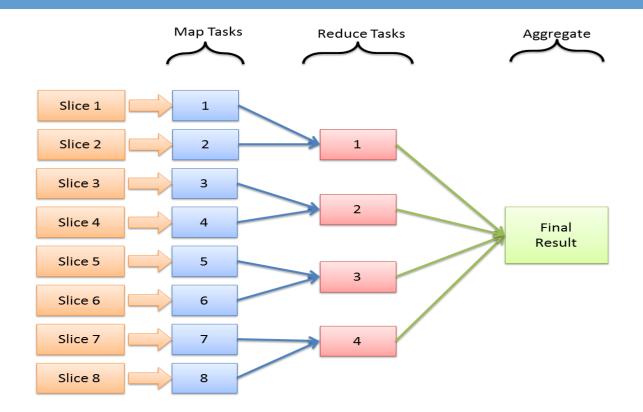




Map >> Reduce





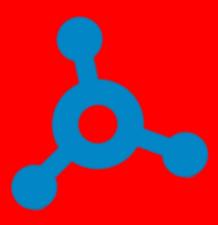


Map >> Reduce

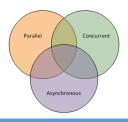


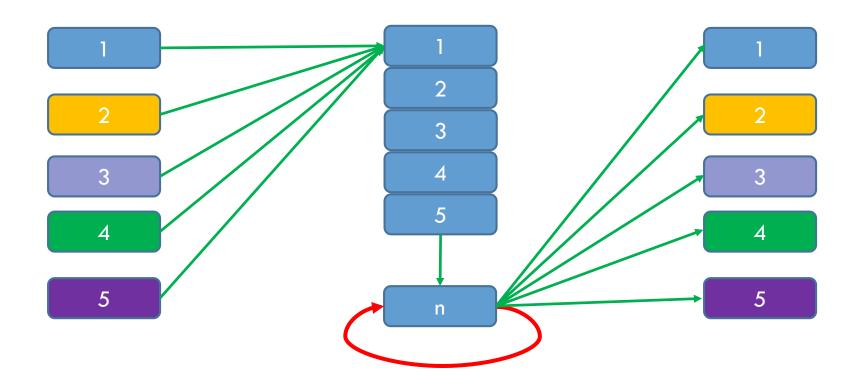
```
var duplicates =
    Directory.GetFiles(path, "*.*", SearchOption.AllDirectories).AsParallel()
        .Select(f =>
            { using (var fs = new FileStream(f, FileMode.Open, FileAccess.Read))
                    return new
                        FileName = f,
                        FileHash = BitConverter.ToString(SHA1.Create().ComputeHash(fs))
                    };
            })
        .GroupBy(f => f.FileHash, EqualityComparer<string>.Default)
        .Select(g => new { FileHash = g.Key, Files = g.Select(z => z.FileName) })
        .Where(g => g.Files.Count > 1)
        //.Skip(1).SelectMany(f => f.Files)
        .ToList();
```

CHANNEL



Channel





```
match msg with
             Read ok when writers.Count = 0 ->
               readers.Enqueue ok
               return! loop()
             Read ok -> let value, cont = writers.Dequeue()
                         Pool.Spawn cont
                         ok value
                         return! loop()
             Write (x,ok) when readers.Count = 0 ->
                     writers.Enqueue(x, ok)
                     return! loop()
             Write (x,ok) -> let cont = readers.Dequeue()
                               Pool.Spawn ok
                               cont x
                               return! loop() }
       loop() )
member this.Read(read) = channelAgent.Post (Read read)
member this.Write(v, ok) = channelAgent.Post (Write (v,ok))
member inline this.Read() = Async.FromContinuations(fun (ok, _, _) -> this.Read ok)
member inline this.Write x = Async.FromContinuations(fun (ok, , ) -> this.Write(x,ok))
```

type [<Sealed>] Channel<'T>() =
 let channelAgent =

let readers = Queue()
let writers = Queue()
let rec loop () = async {

MailboxProcessor<ChannelMsg<'T>>.Start(fun inbox ->

let! msg = inbox.Receive()

atomic { stm }

Software Transactional Memory

```
val readTVar : TVar<'a> -> Stm<'a>
                  val writeTVar : TVar<'a> -> 'a -> Stm<unit>
                  val: atomically: Stm<'a>-> 'a
                  let incr x =
                                                                   atomic increment
                     stm {
                       let! v = readTVar x
                       let! = writeTVar x (v+1)
introduce
                       return v
atomic
block
                  let incr2x =
                     stm {
                                                       composable transactions
                       let! = incr x
                       let! v = incr x
                       return v
```

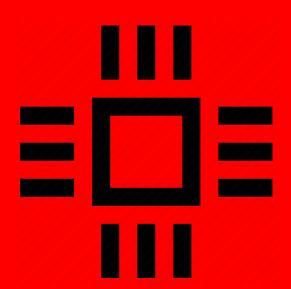
Software Transactional Memory

```
let take fork = stm {
                      let! ref fork = readTVar (fork)
                      do! check ref fork
                      do! writeTVar (fork) false }
let acquire left_fork right_fork =
             stm {
                 do! take left fork
                 do! take right fork } |> atomically
let release left fork right fork =
    stm {
        do! writeTVar left fork true
        do! writeTVar right fork true
    } |> atomically
```

Software Transactional Memory

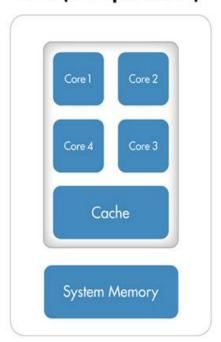


GP-GPU

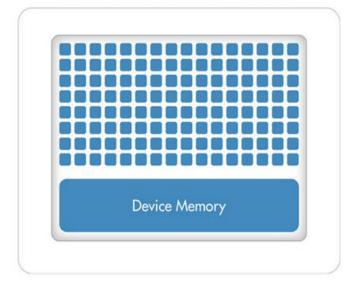


GPU

CPU (Multiple Cores)



GPU (Hundreds of Cores)

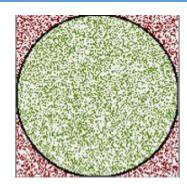




GPU - MSFT Accelerator

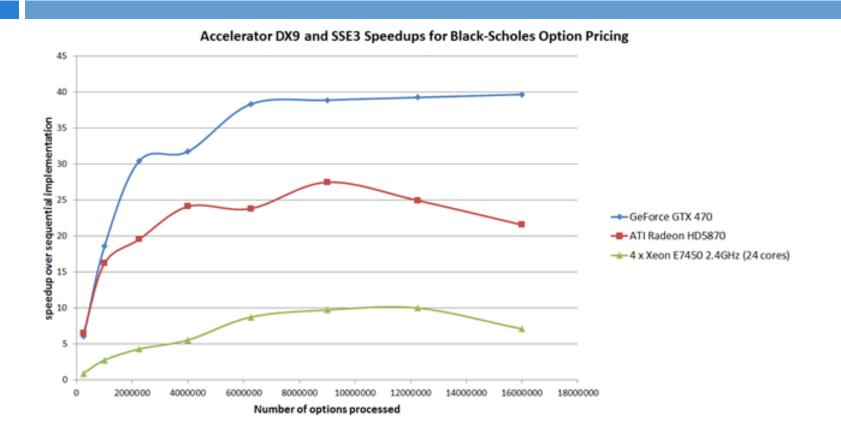
□ Microsoft Accelerator

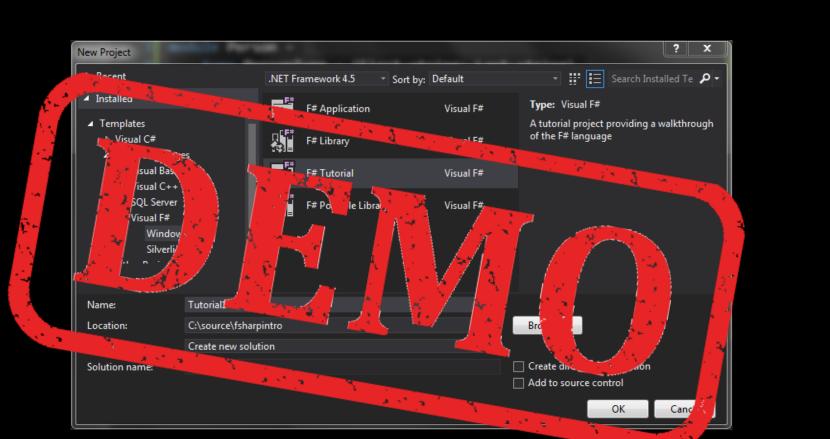
The project Microsoft Research Accelerator is a .Net library for developing array-based computations and executing them in parallel on multicore CPUs or more interestingly, using GPU shaders



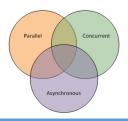
 Accelerator handles all the details of parallelizing and running the computation on the selected target processor, including GPUs and multicore CPUs

GPU - MSFT Accelerator





Summary



Embrace Parallelism exploiting the potential of multicore CPUs

No single solution...

You must combine different programming models

(right tool for the job)

FP really **shines** in the area of concurrency





The tools we use have a profound (and devious!) influence on our thinking habits, and, therefore, on our thinking abilities.

-- Edsger Dijkstra

How to reach me





github.com/rikace/Presentations/AkkaActorModel

meetup.com/DC-fsharp

@DCFsharp

@TRikace

tericcardo@gmail.com

https://www.surveymonkey.com/r/S6PV89W