



Taming and Composing High-Performance Stream Processing

reactive programming is a programming paradigm that maintains a continuous interaction with their environment, but at a speed which is determined by the environment, not the program itself

Gérard Berry

Agenda

The rise of Reactive Streams

Reactive Extensions and back pressure

Actor Model, does not fit well

Reactive Streams

Akka Streams

Short demo & code review

Objectives

- React to high rate events
- Compose events from different sources
- Simplify complex logic
- Pipeline for combinatorics over event stream

Introduction - Riccardo Terrell

- ④ Originally from Italy, currently - Living/working in Washington DC ~10 years
- ④ +/- 19 years in professional programming
 - ④ C++/VB → Java → .Net C# → Scala → Haskell → C# & F# → ??
- ④ Author of the book “Functional Concurrency in .NET” - Manning
- ④ *Polyglot programmer - believes in the art of finding the right tool for the job*



@trikace

rickyterrell.com

tericcardo@gmail.com



Why Reactive?

- Billions of internet connected devices
- Data is transformed and pushed continuously
- Services must be always up
- Must be fast

GOOGLE

- Current storage = 15 exabytes
- Processed per day = 100 petabytes
- Number of pages indexed = 60 trillion
- Unique search users per month > 1 billion
- Searches per second = 2.3 million

one exabyte is one
quintillion bytes

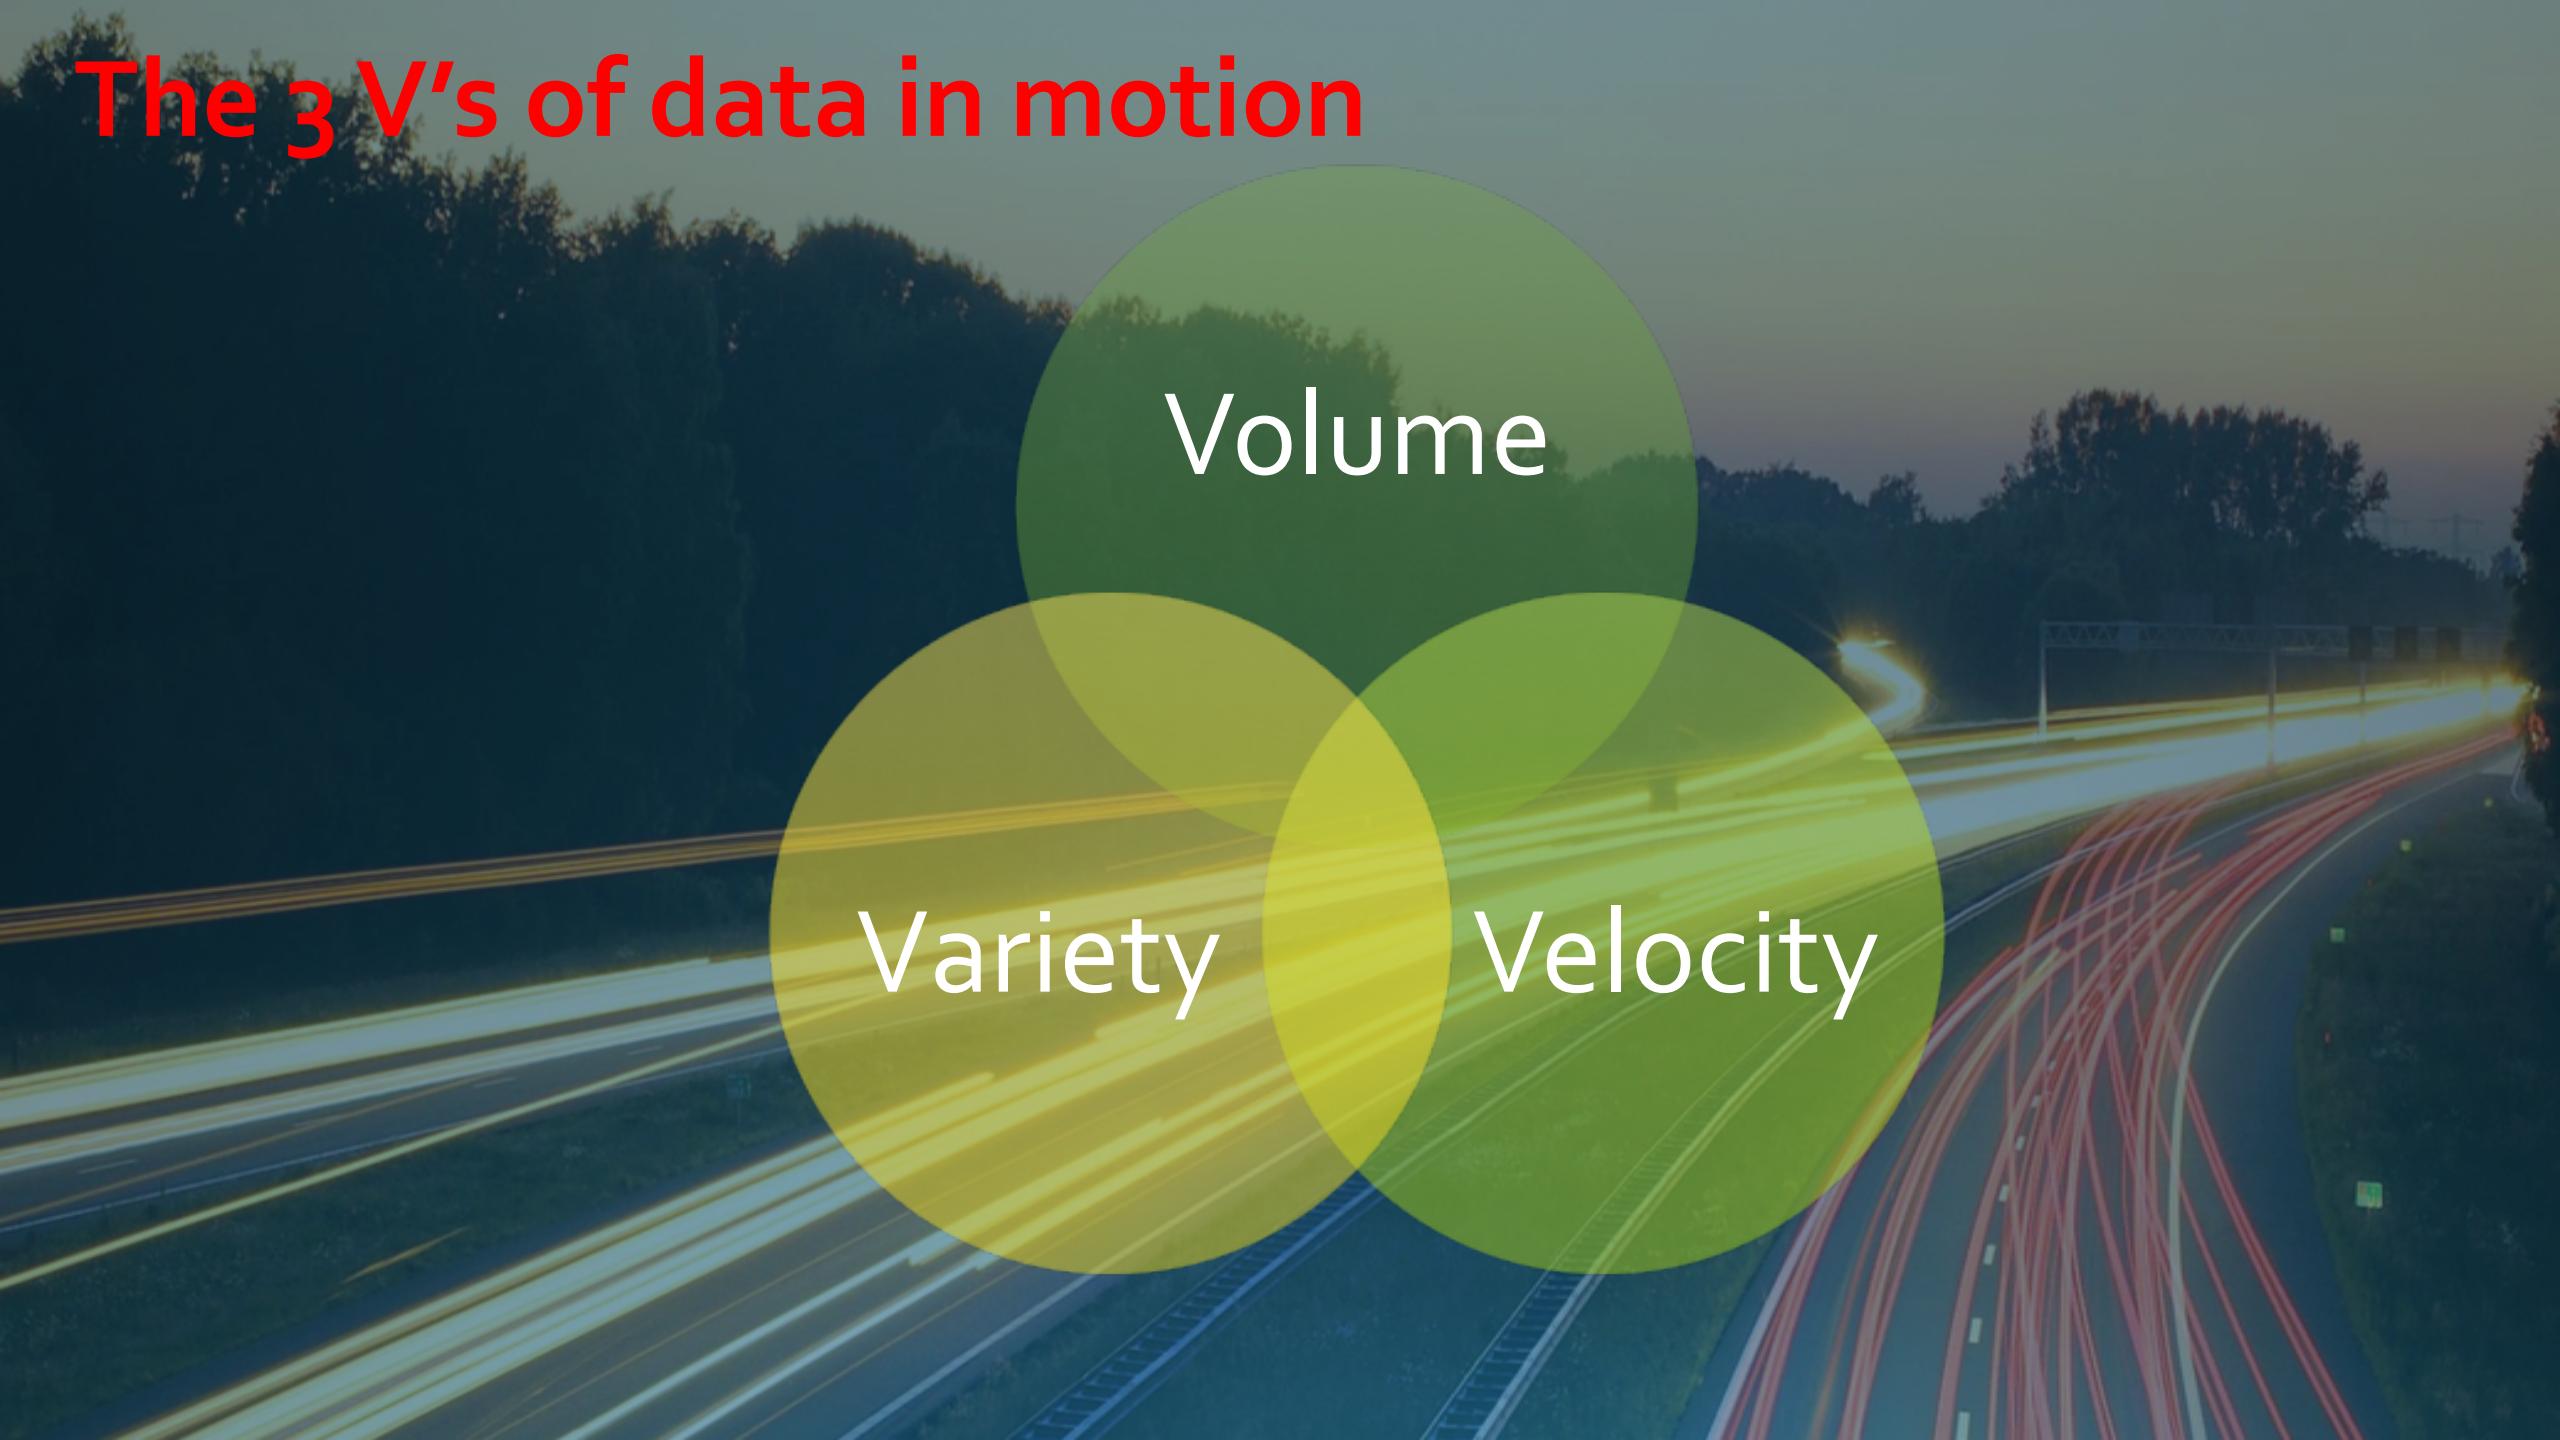
bytes = 1000000000000000000
1000 petabytes
1 million terabytes
1 billion gigabytes

The Rise of Event Stream Systems

It's all about data these days. Data helps us make informed decisions. Big Data helps us make informed and insightful decisions. Big streams of data help us make informed, insightful and timely decisions. These continuously flowing streams of data are often called event streams. It's increasingly common to build software systems whose primary purpose is to process event streams.



The 3 V's of data in motion



Volume

Variety

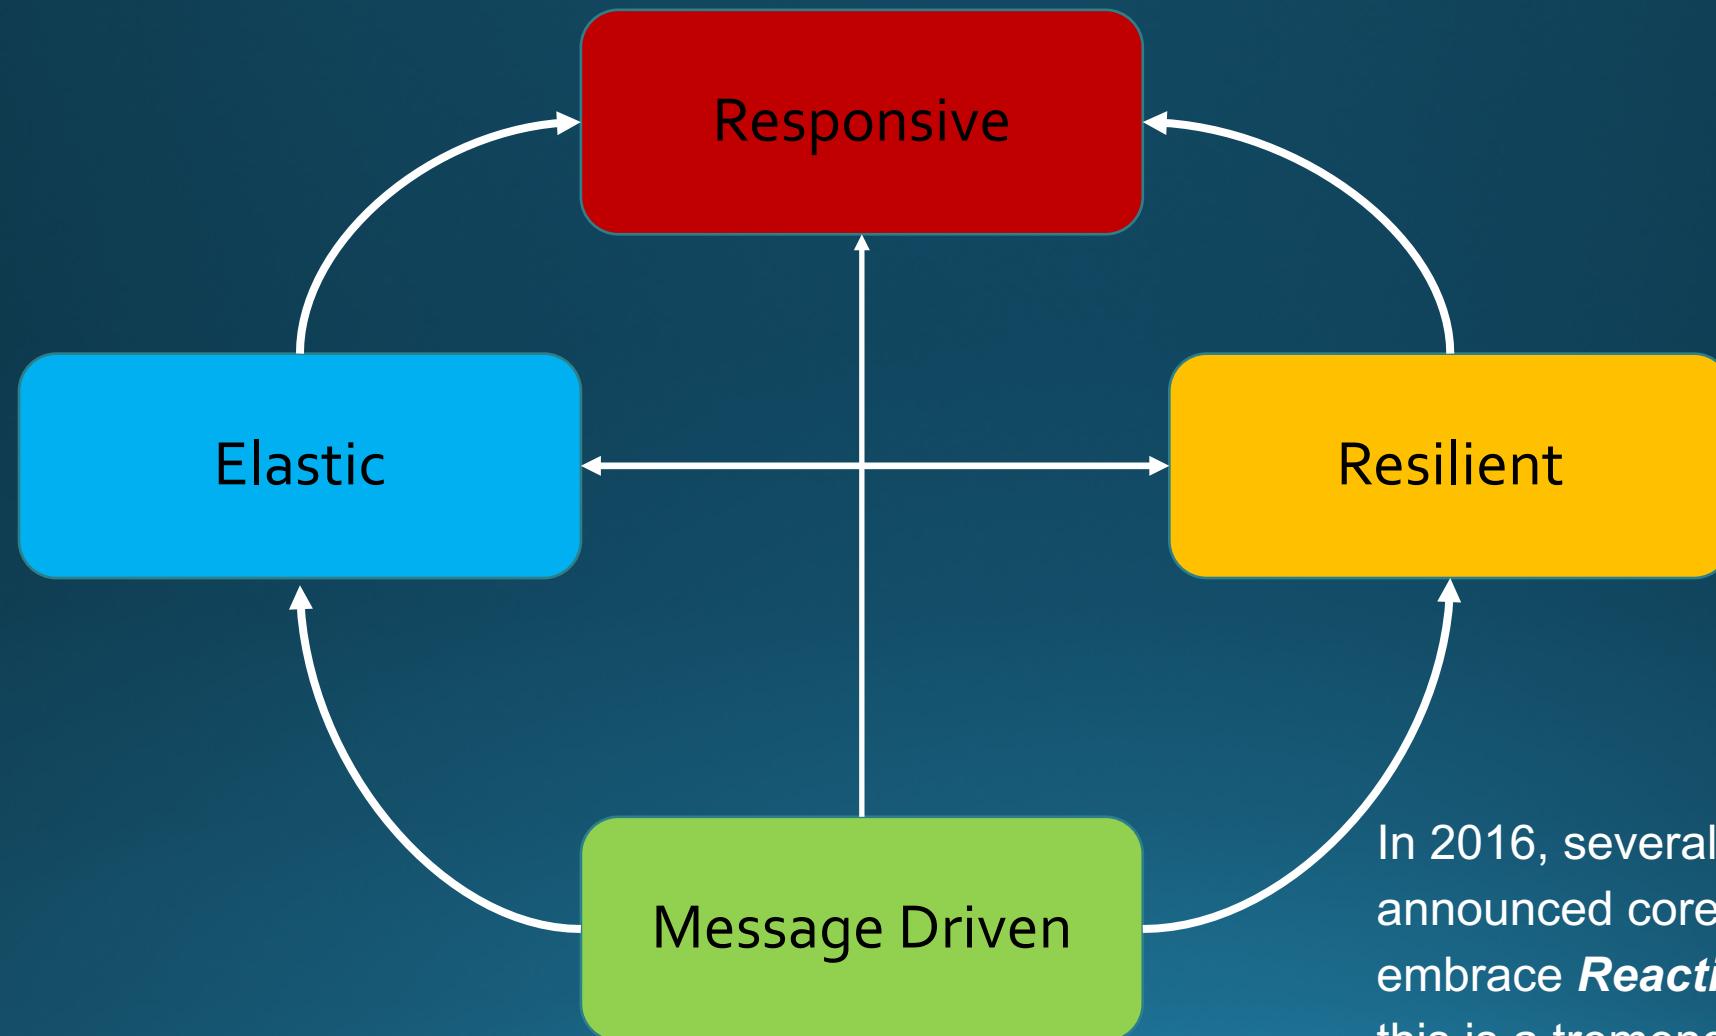
Velocity

New Tools for a New Era

We need to build systems that:

- *react* to events — Event-Driven
- *react* to load — Scalable
- *react* to failure — Resilient
- *react* to users — Responsive

Reactive Manifesto



In 2016, several major vendors have announced core initiatives to embrace ***Reactive Programming*** this is a tremendous validation of the problems faced by companies today.

Functional Reactive Programming

Functional Reactive Programming

- declarative
- functions as values
- side-effects free
- referential transparency
- immutable
- composition



Twitter Paper

Your Server as a Function

Marius Eriksen

Twitter Inc.

marius@twitter.com

Abstract

Building server software in a large-scale setting, where systems exhibit a high degree of concurrency and environmental variability, is a challenging task to even the most experienced programmer. Efficiency, safety, and robustness are paramount—goals which have traditionally conflicted with modularity, reusability, and flexibility.

We describe three abstractions which combine to present a powerful interface.

Services Systems boundaries are represented by asynchronous functions called *services*. They provide a symmetric and uniform API: the same abstraction represents both clients and servers.

Filters Application-agnostic concerns (e.g. timeouts, retries, authentication) are encapsulated by *filters* which compose to build services from multiple independent modules.



Twitter Paper

Your Server as a Function

Servers in a large-scale setting are required to process tens of thousands, if not hundreds of thousands of requests concurrently; they need to handle partial failures, adapt to changes in network conditions, and be tolerant of operator errors. As if that weren't enough, harnessing off-the-shelf software requires interfacing with a heterogeneous set of components, each designed for a different purpose. These goals are often at odds with creating modular and reusable software [6].

Abstract

Building server software exhibits a high degree of concurrency, a challenging task to ensure efficiency, safety, and robustness. Traditionally, these goals have traditionally conflicted with one another, as they are often represented by asynchronous operations that provide a symmetric and uniform interface for both clients and servers.

We present three abstractions around which we structure our server software at Twitter. They adhere to the style of functional programming—emphasizing immutability, the composition of first-class functions, and the isolation of side effects—and combine to present a large gain in flexibility, simplicity, ease of reasoning, and robustness.

represented by asynchronous operations that provide a symmetric and uniform interface for both clients and servers.

(e.g. timeouts, retries, automatic failover) which compose to build complex systems.

It's really clear that the imperative style of programming has run its course. ... We're sort of done with that. ... However, in the declarative realm we can speculate a 10x improvement in productivity in certain domains.

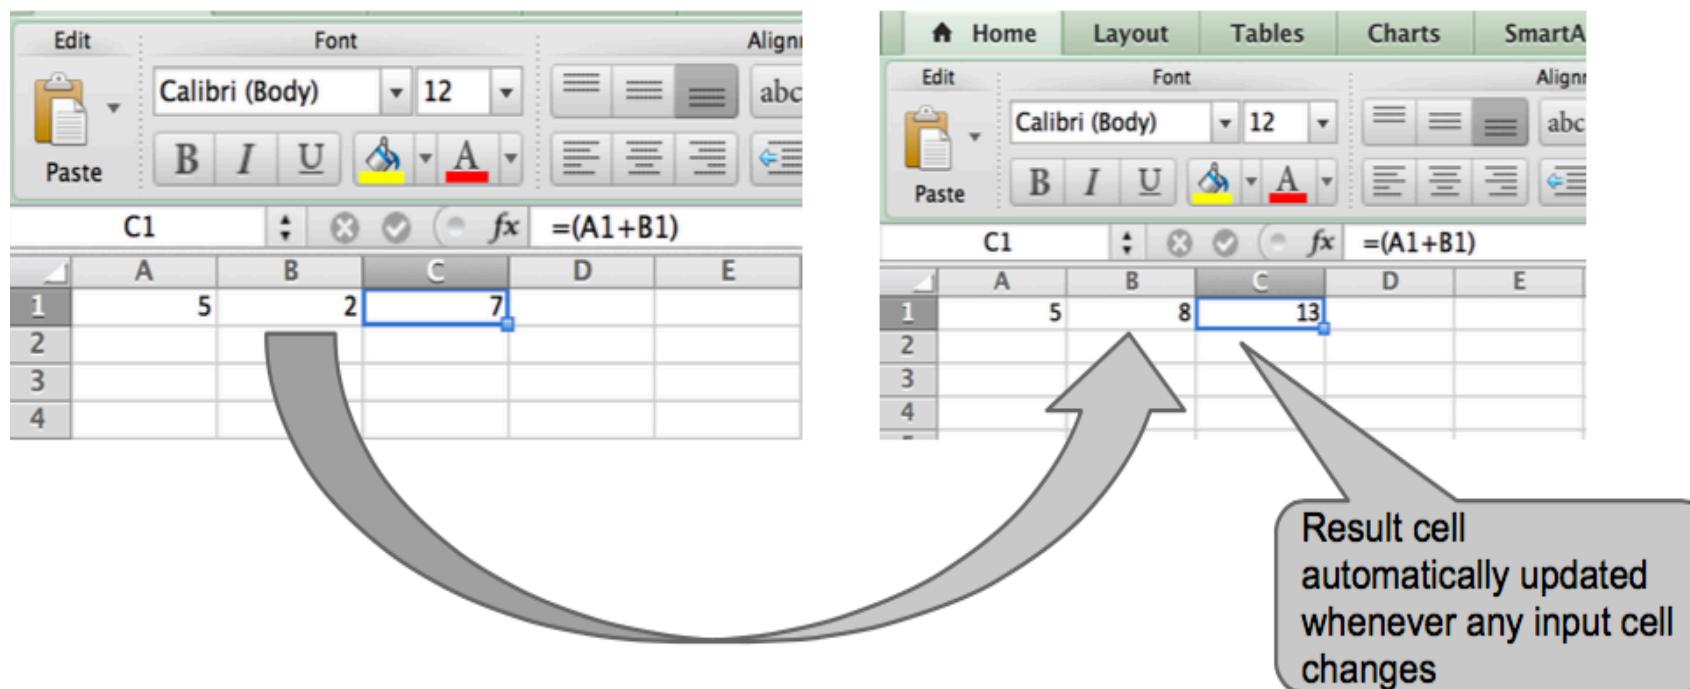
-Anders Hejlsberg
C# Architect

(from his MIX07 keynote)

Functional Reactive Programming

What is Reactive Programming

**SpreadSheet == Mother of All
Reactive Programming**



Functional Reactive Programming

Push-Pull Functional Reactive Programming

Push-Pull Functional Reactive Programming

Conal Elliott

LambdaPix

conal@conal.net

Abstract

Functional reactive programming (FRP) has simple and powerful semantics, but has resisted efficient implementation. In particular, most past implementations have used demand-driven sampling, which accommodates FRP's continuous time semantics and fits well with the nature of functional programming. Consequently, values are wastefully recomputed even when inputs don't change.

more composable than their finite counterparts, because they can be scaled arbitrarily in time or space, before being clipped to a finite time/space window.

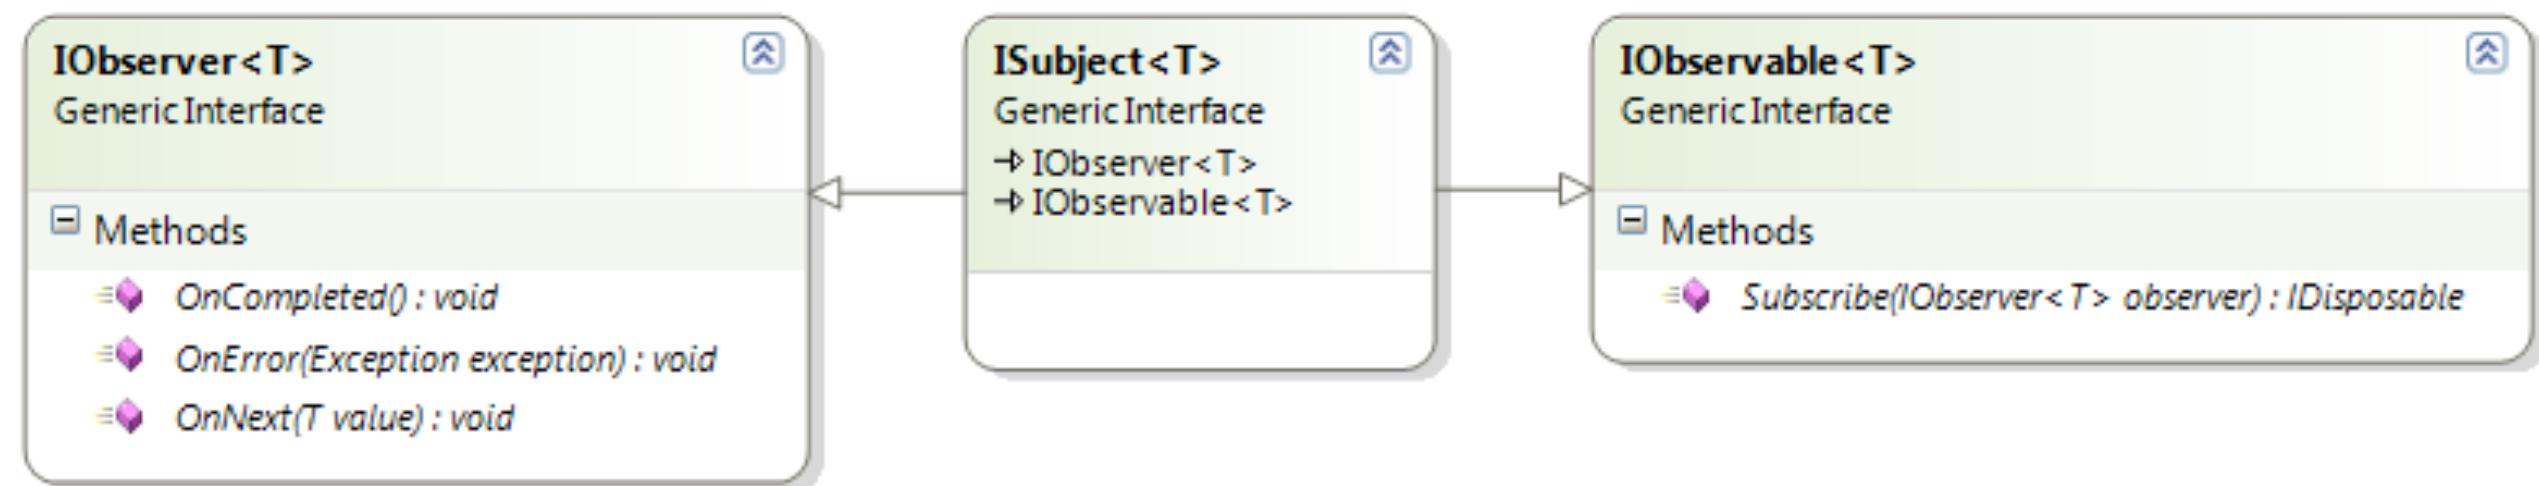
While FRP has simple, pure, and composable semantics, its efficient implementation has not been so simple. In particular, past implementations have used demand-driven (pull) sampling of reactive behaviors, in contrast to the data-driven (push) evaluation typ-

The logo for Reactive Extensions features a circular emblem composed of three interlocking, rounded shapes in shades of pink and purple. The emblem is centered on a dark teal background. Overlaid on the bottom half of the emblem is the text "Reactive Extensions" in a large, white, sans-serif font.

Reactive Extensions

What are Reactive Extensions

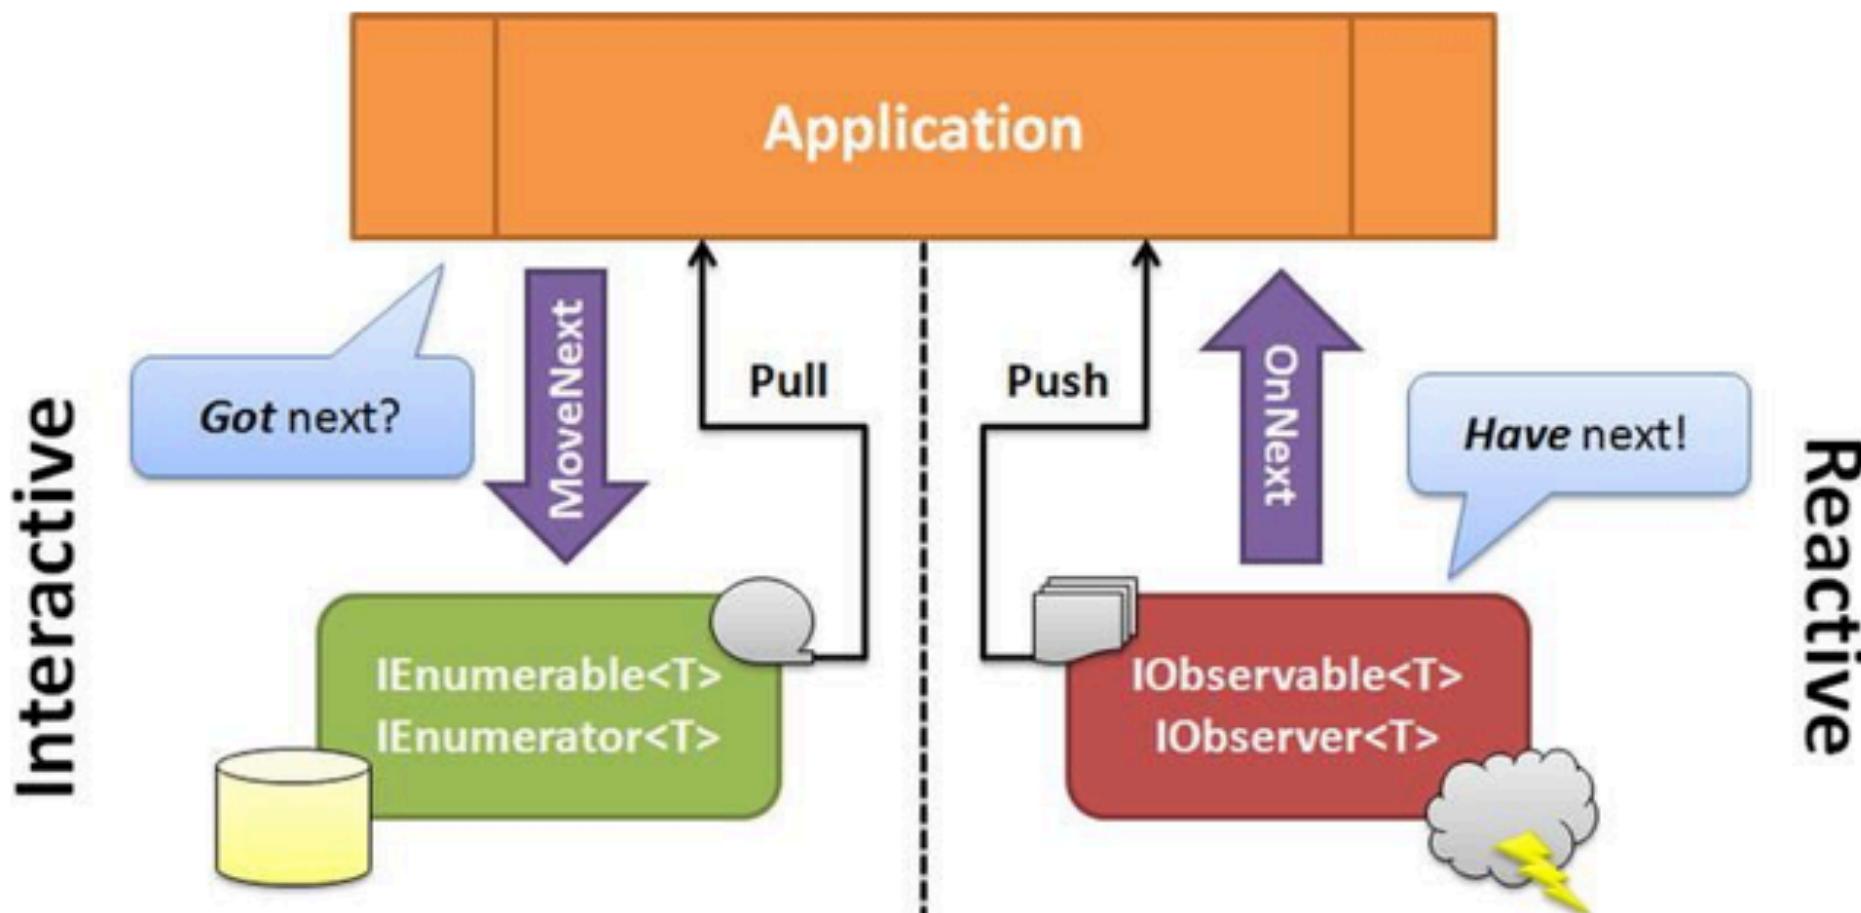
Rx is a library for composing asynchronous and event-based programs using observable sequences and LINQ-style query operators



Push vs Pull



Push vs Pull



IObserver & IObservable

```
type IObserver<'a> = interface
    abstract OnCompleted : unit -> unit
    abstract OnError : exn -> unit
    abstract OnNext : 'a -> unit
end
```

```
type IObservable<'a> = interface
    abstract Subscribe : IObserver<'a> -> IDisposable
end
```



I`Observable` the dual of I`Enumerable`

```
type IObserver<'a> = interface
    abstract OnNext : 'a with set
    abstract OnCompleted : unit -> unit
    abstract OnError : Exception -> unit
end
```

```
type IObservable<'a> = interface
    abstract Subscribe : IObserver<'a>
        -> IDisposable
end
```



```
type IEnumerator<'a> = interface
    interface IDisposable
    interface IEnumerator
end

abstract Current : 'a with get
abstract MoveNext : unit -> bool
end

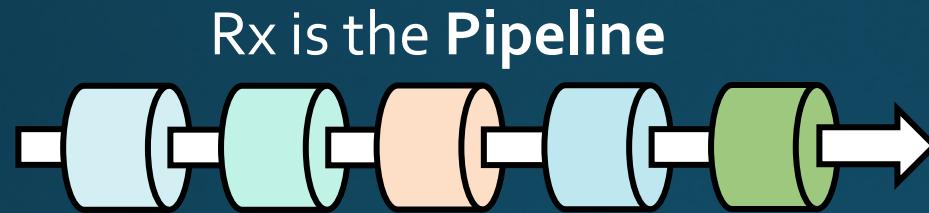
type IEnumerable<'a> = interface
    interface IEnumerable
end

abstract GetEnumerator : IEnumerator<'a>
end
```

What is Rx?

Rx **compose pipeline** of operations over **single or multiple event's source**

Focus on what happens **between** the Producer and the Consumer

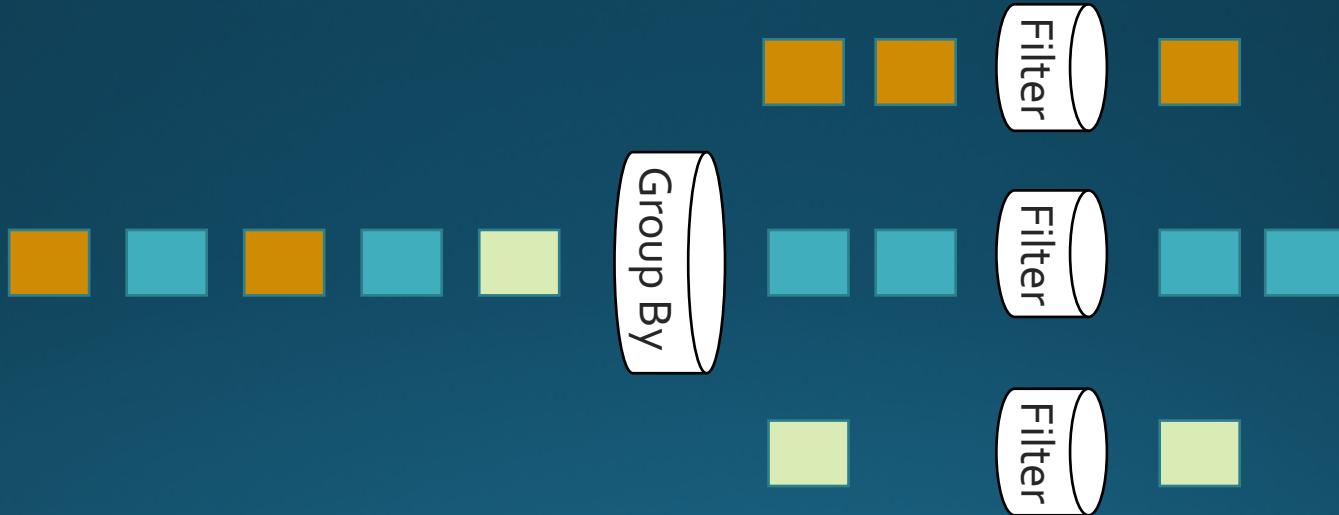


Rx = Async Data Stream Composition

Rx = LINQ to Async Data Stream

What is Rx?

Rx **compose pipeline** of operations over **single or multiple event's source**
Focus on what happens **between** the Producer and the Consumer

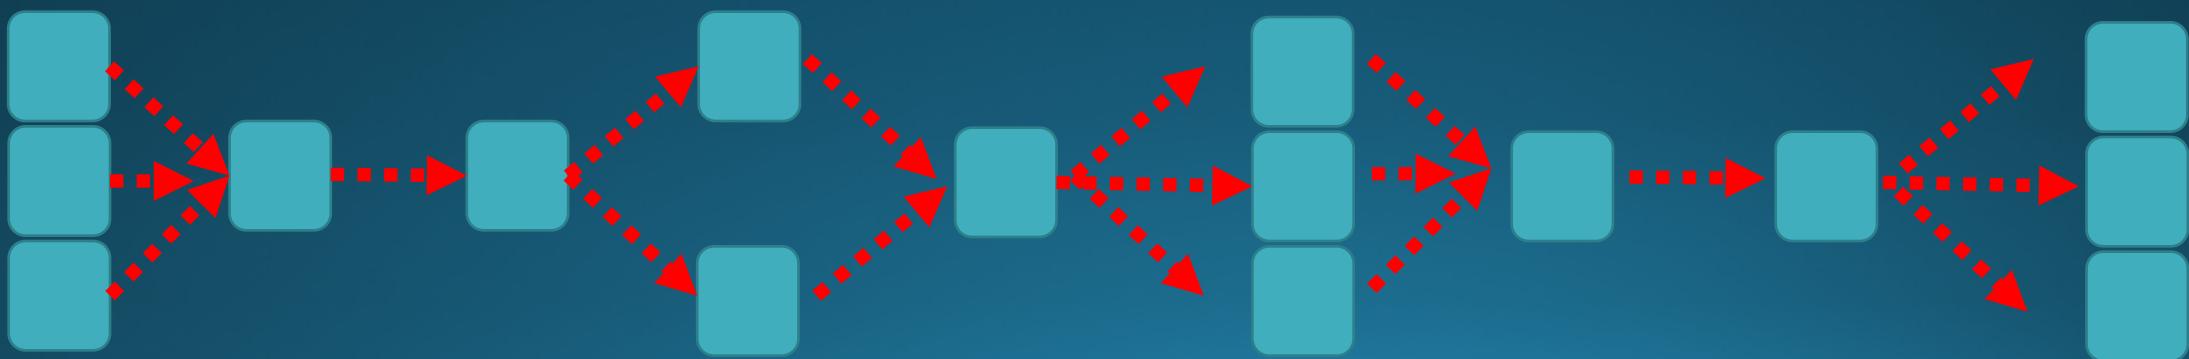


Rx = Async Data Stream Composition

Rx = LINQ to Async Data Stream

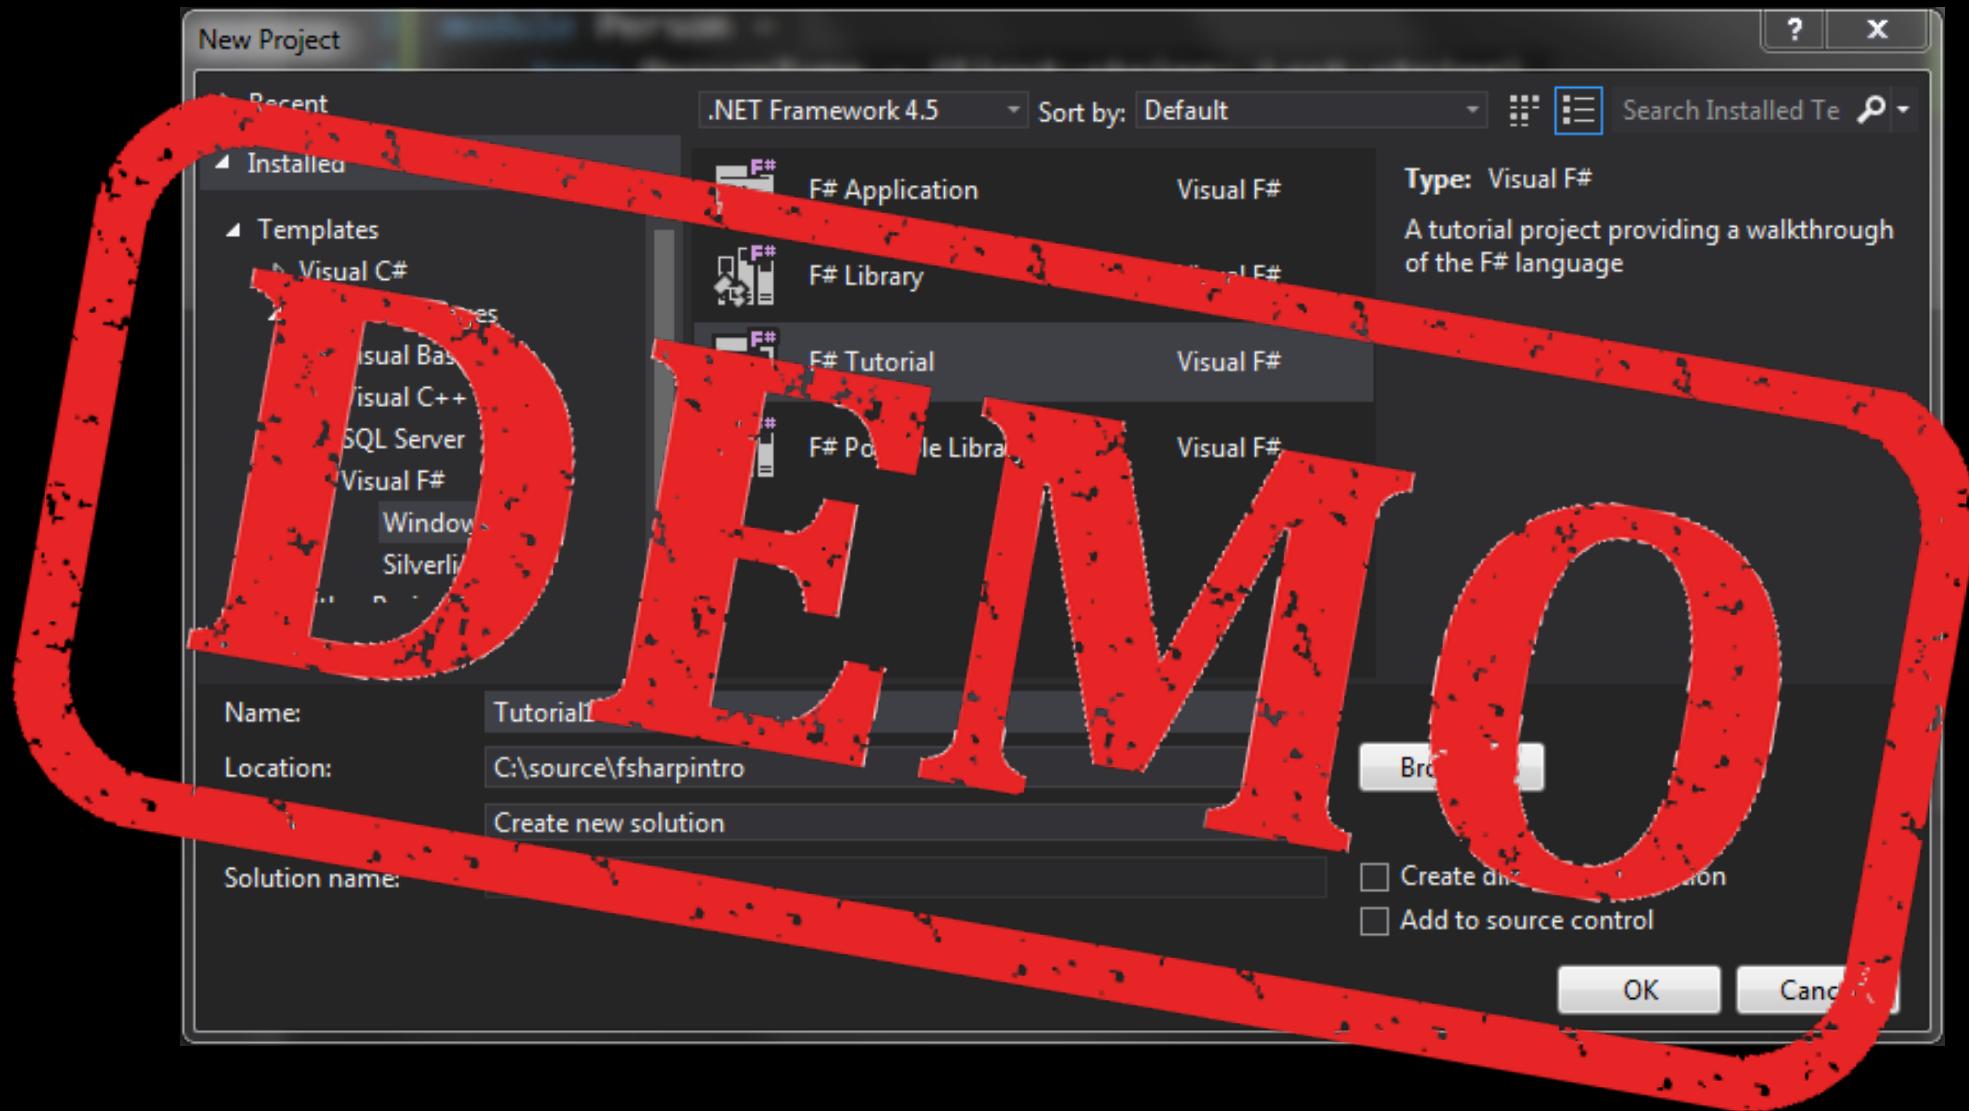
What is a Stream?

- ephemeral flow of data
- possibly unbounded in size
- focused on describing transformation
- can be formed into processing networks



Reactive Extensions

```
stockFiles
    .ObservableStreams(StockData.Parse, 10)
    .Throttle(TimeSpan.FromMilliseconds(10))
    .ObserveOn(ctx)
    .GroupBy(stocks => stocks.Symbol)
    .SelectMany(group =>
        .Subscribe(stock =>
            UpdateChart(chart, stock)));
    
```



Reactive Extensions

```
stockFiles
    .ObservableStreams(StockData.Parse, 10)
    .Throttle(TimeSpan.FromMilliseconds(10))
    .ObserveOn(ctx)
    .GroupBy(stocks => stocks.Symbol)
    .SelectMany(group =>
        .Subscribe(stock =>
            UpdateChart(chart, stock)));

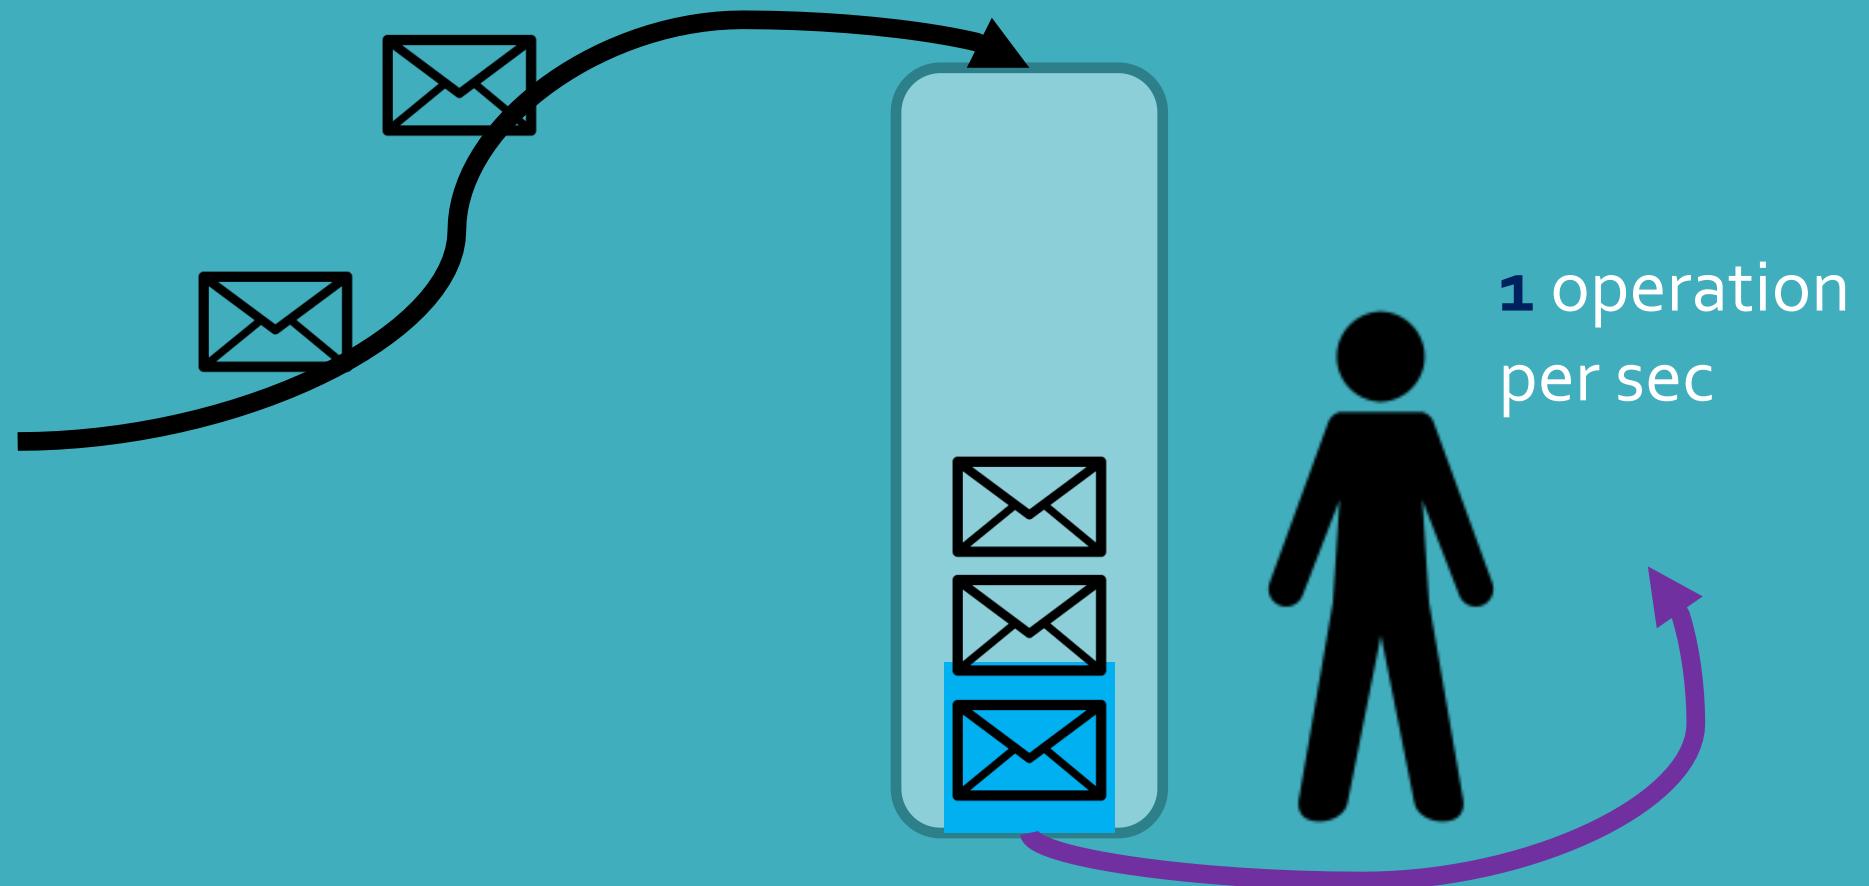
```

Throttle is destructive

Push Model

Subscriber usually has some kind of buffer.

5 operations
per sec

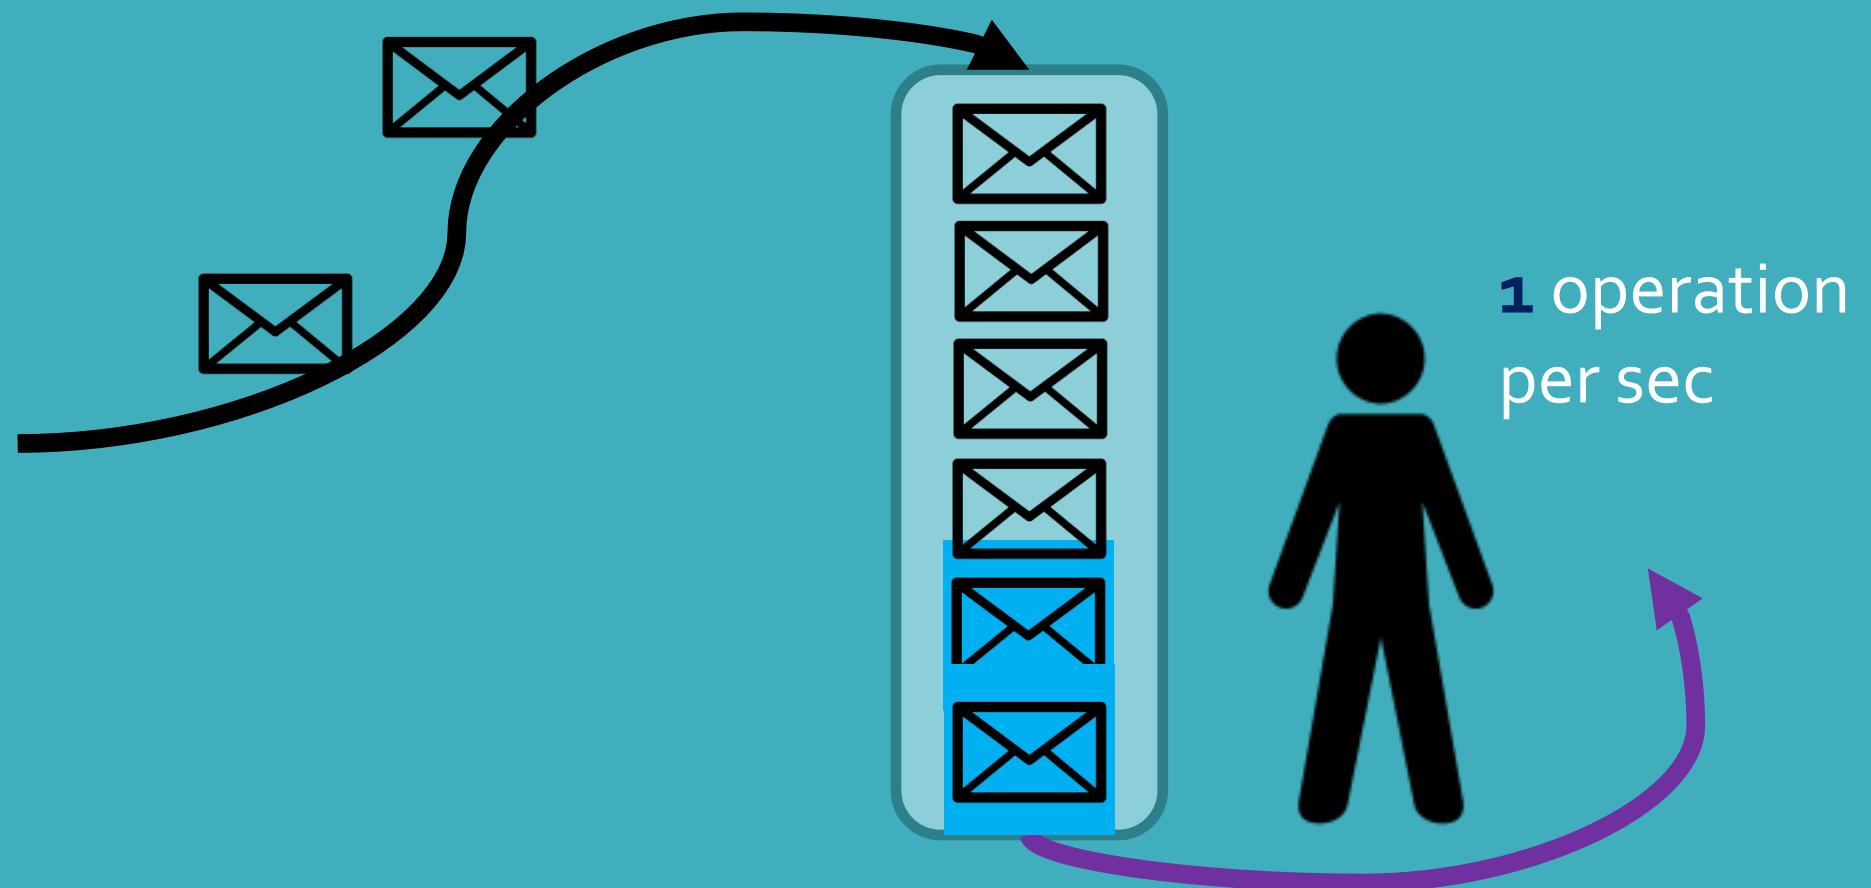


1 operation
per sec

Push Model

Subscriber usually has some kind of buffer.

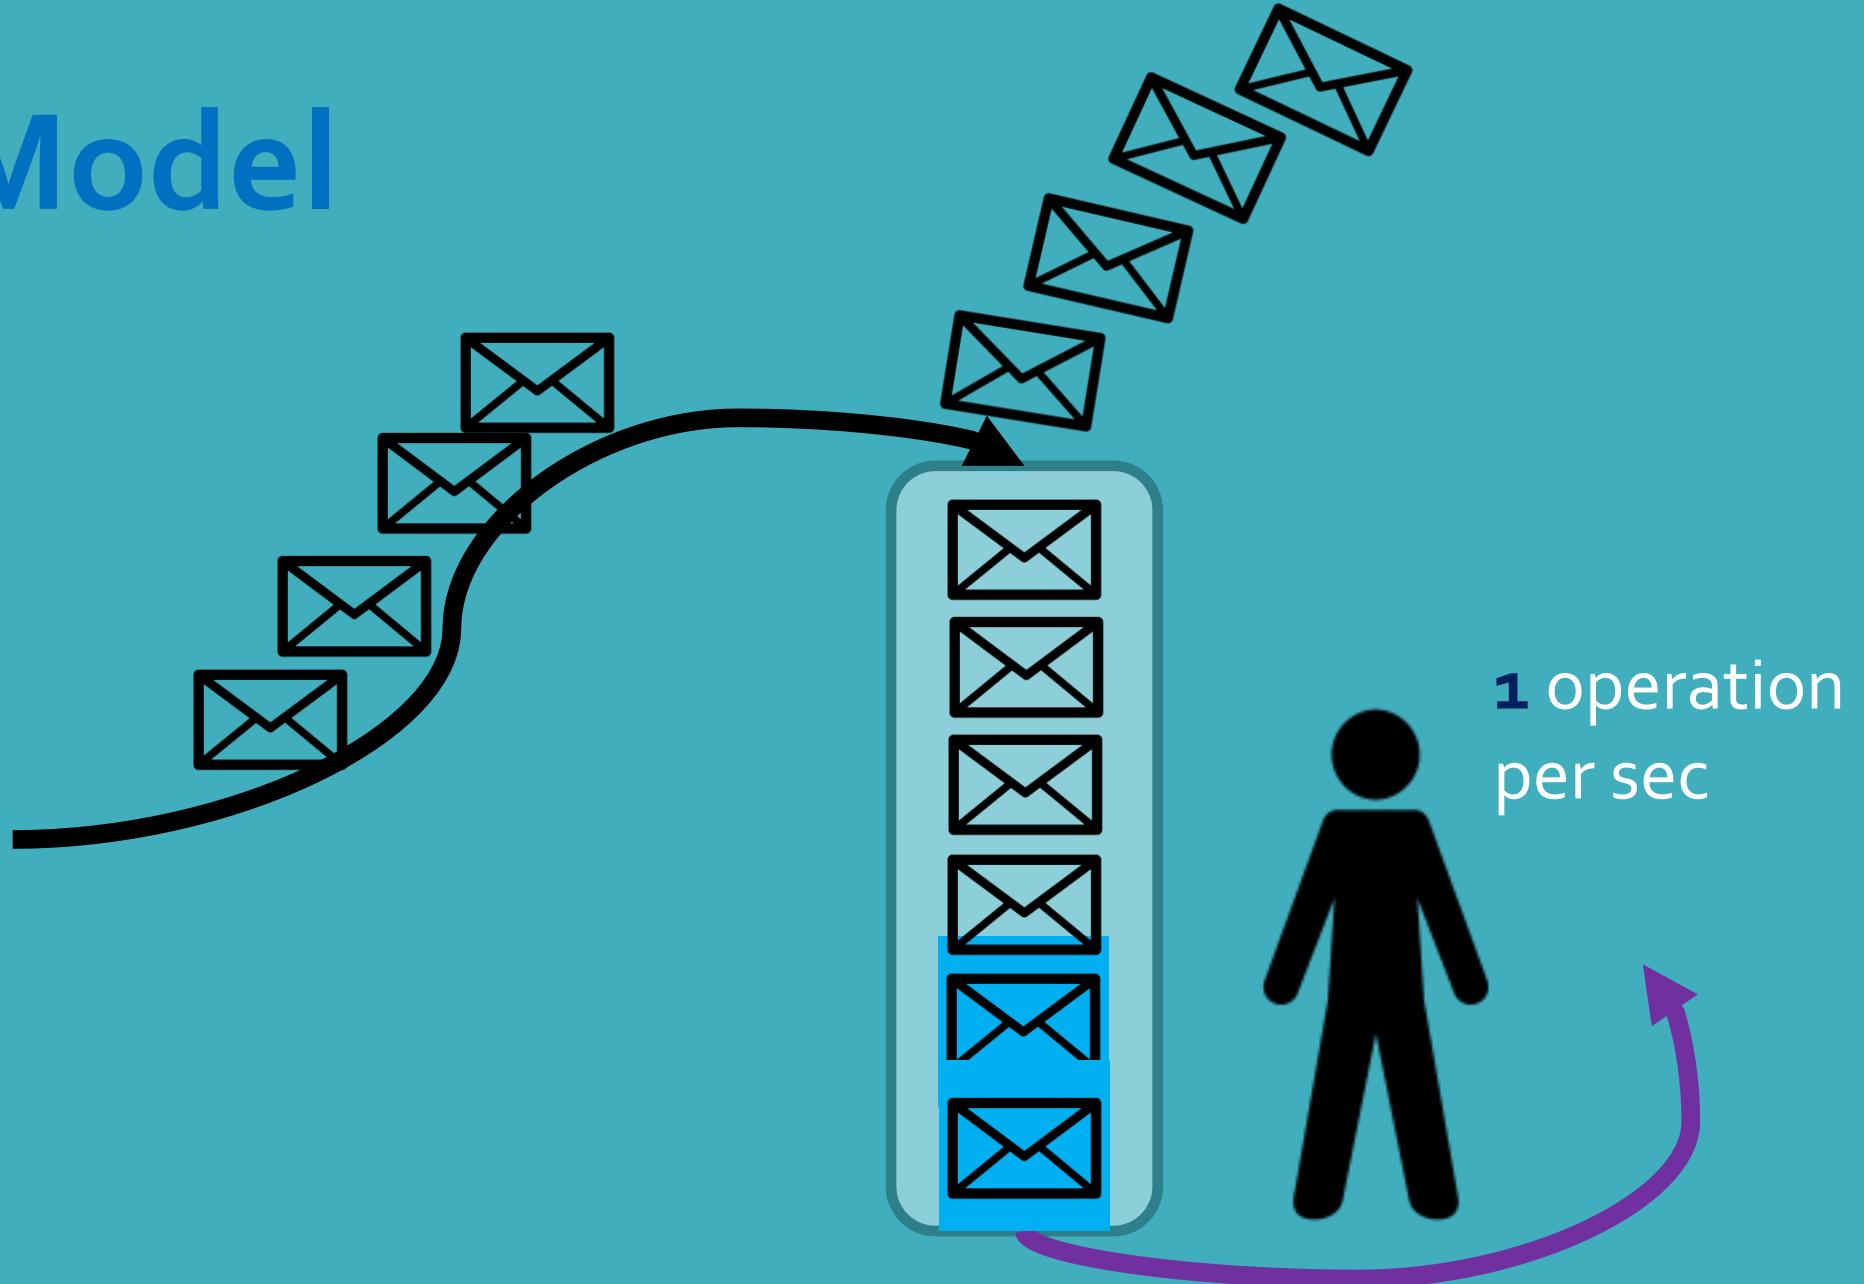
5 operations
per sec



1 operation
per sec

Push Model

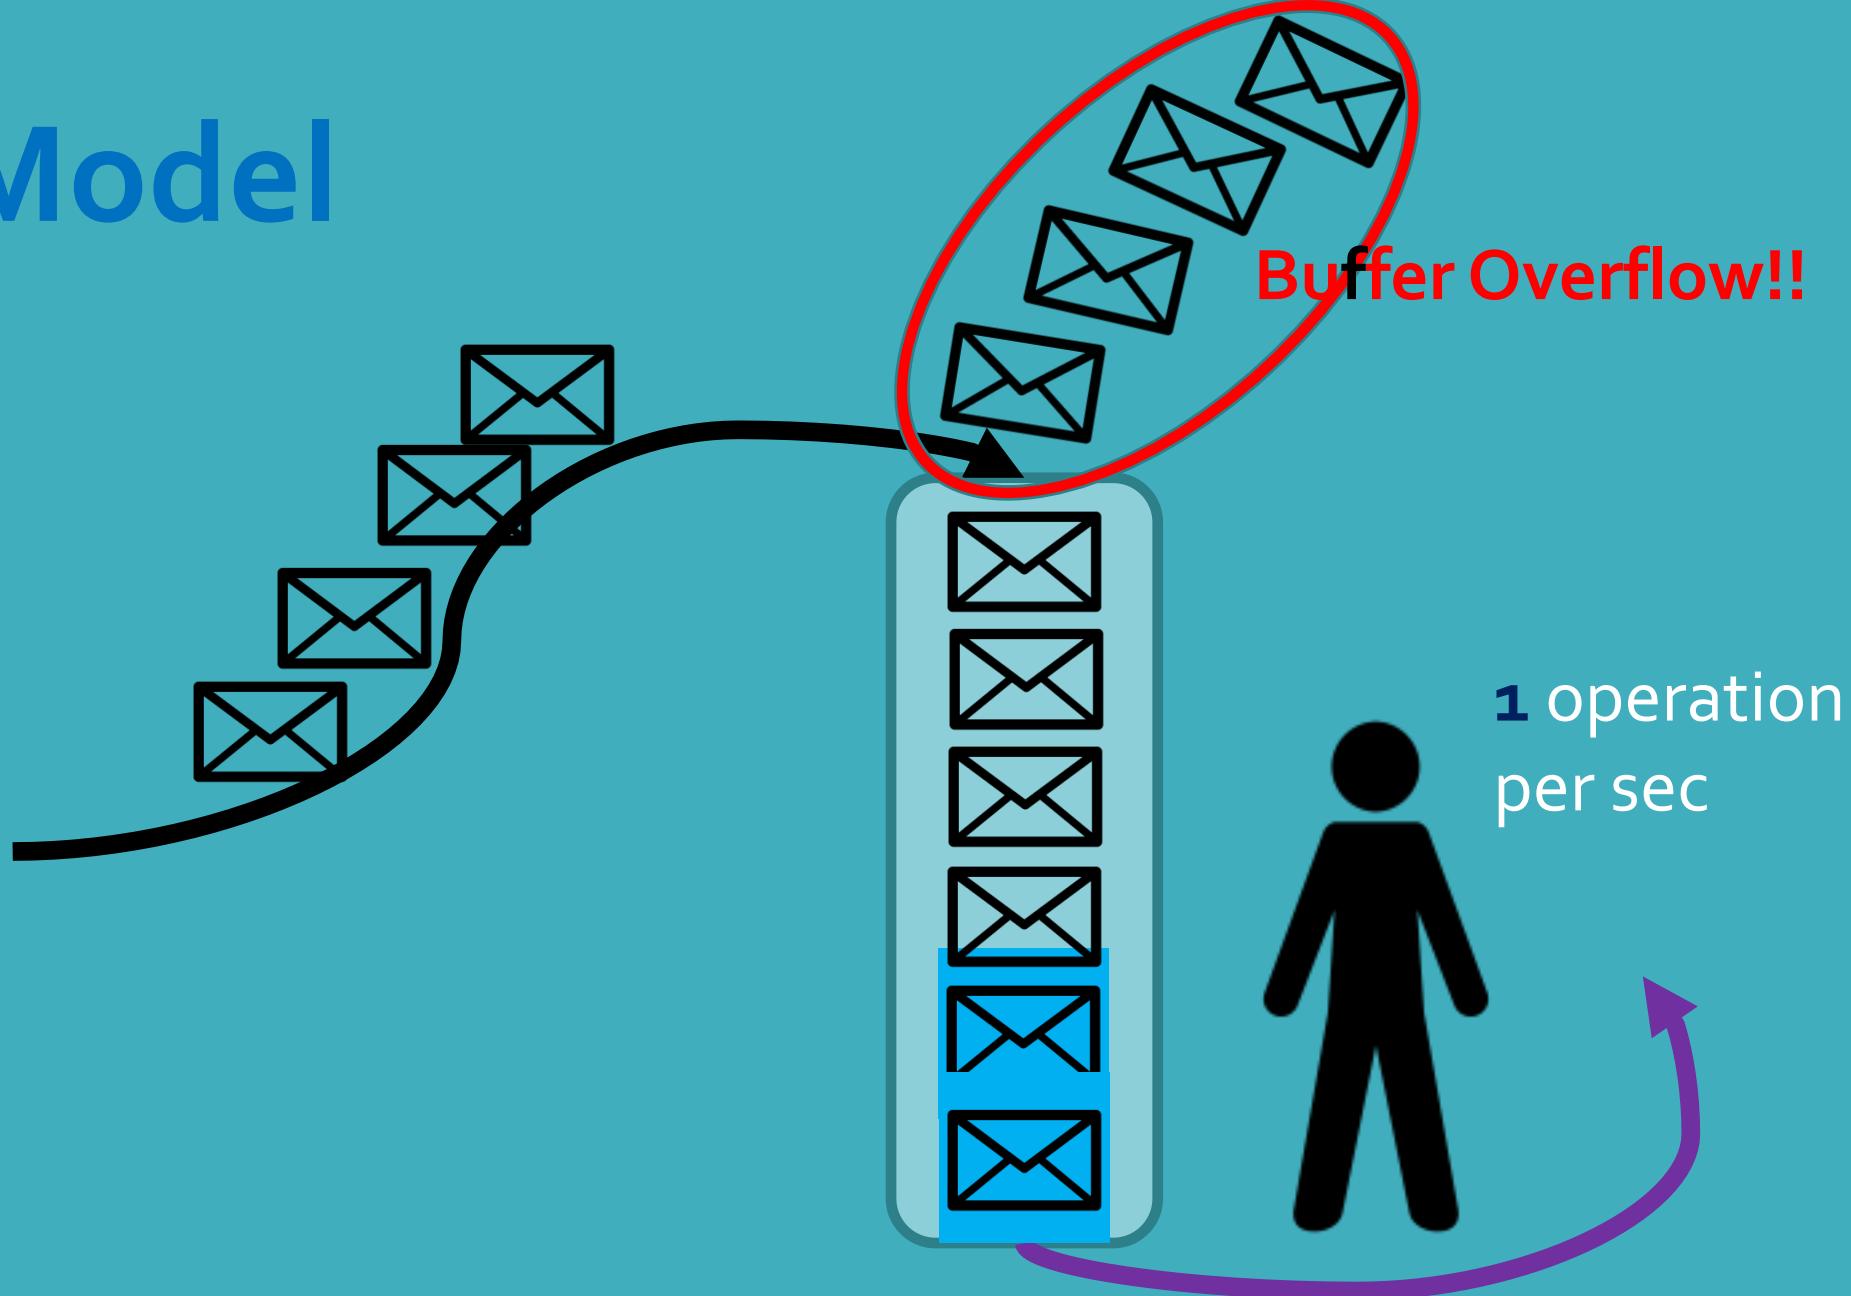
5 operations
per sec



1 operation
per sec

Push Model

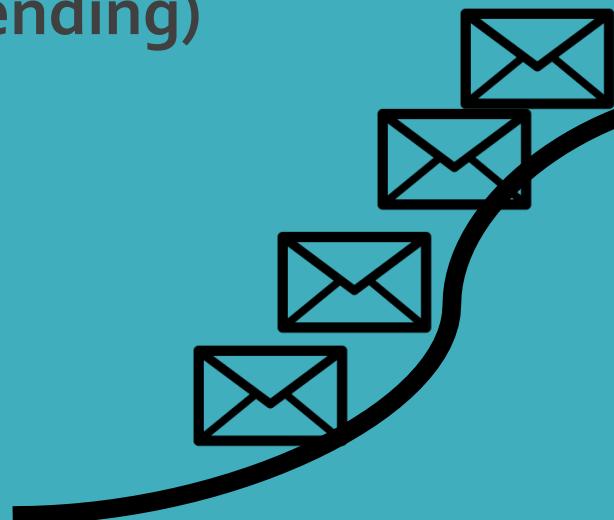
5 operations
per sec



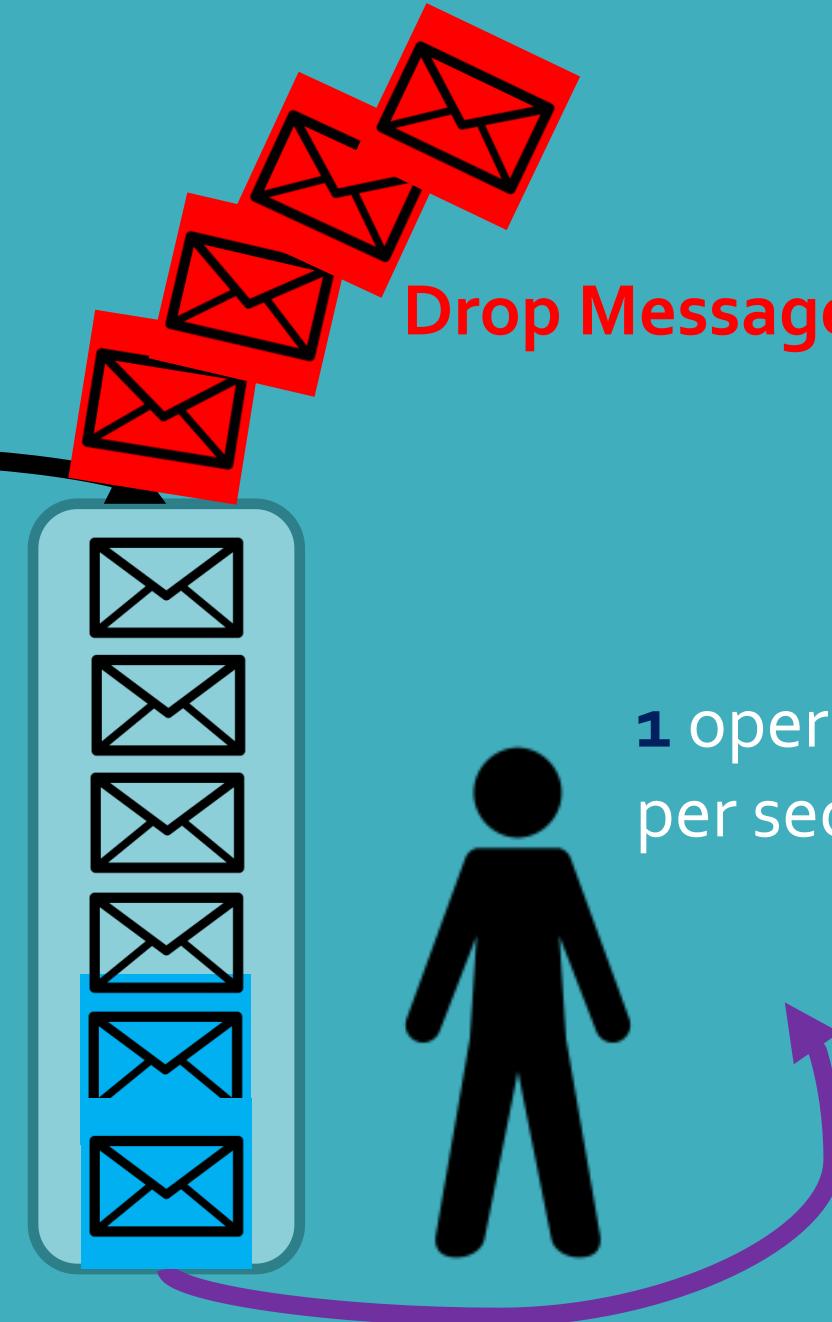
Push Model

Bounded buffer drop exceeded messages
(require re-sending)

5 operations
per sec



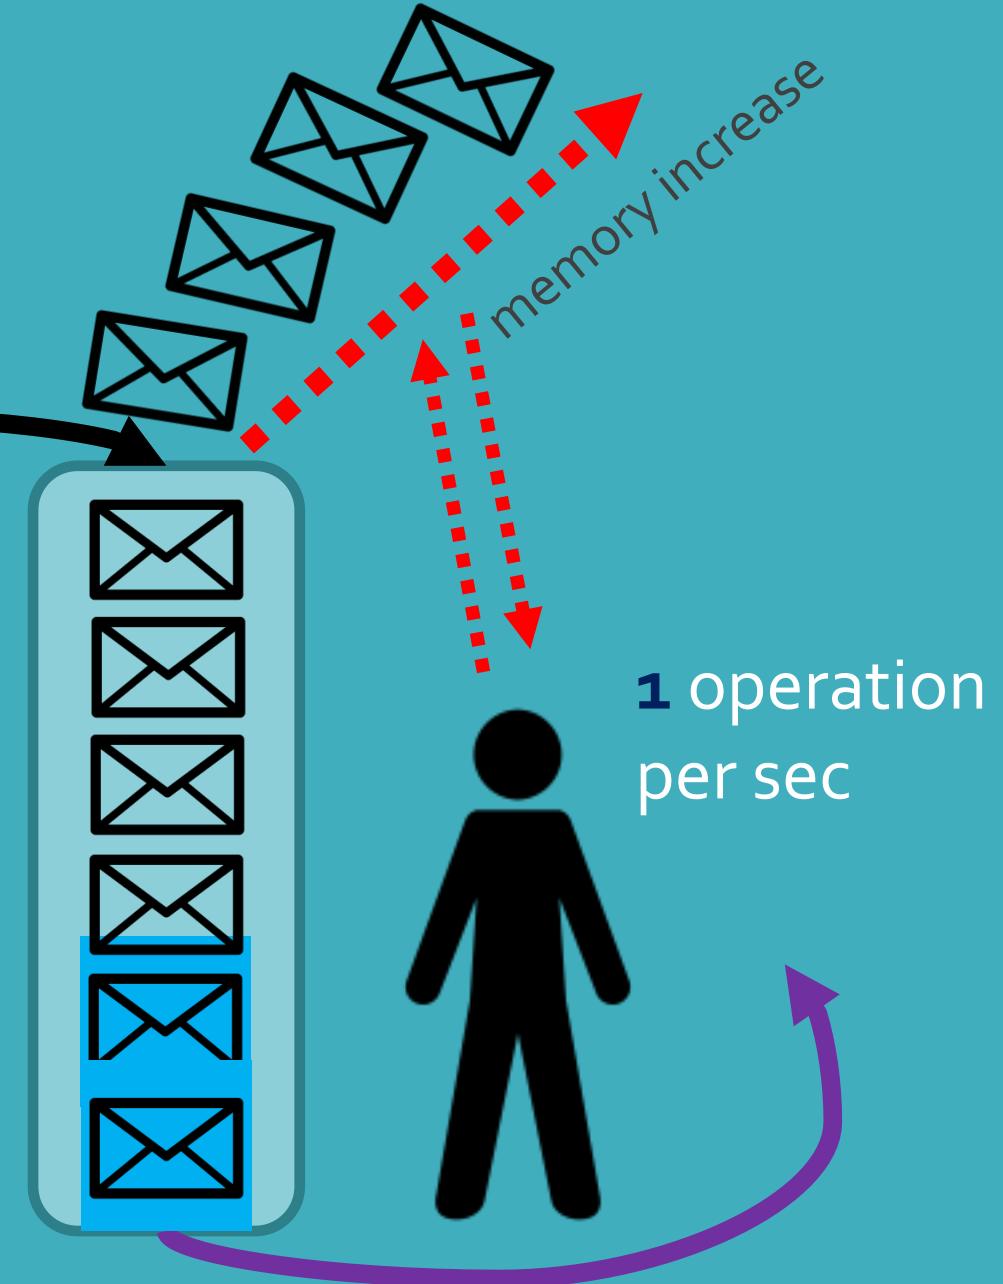
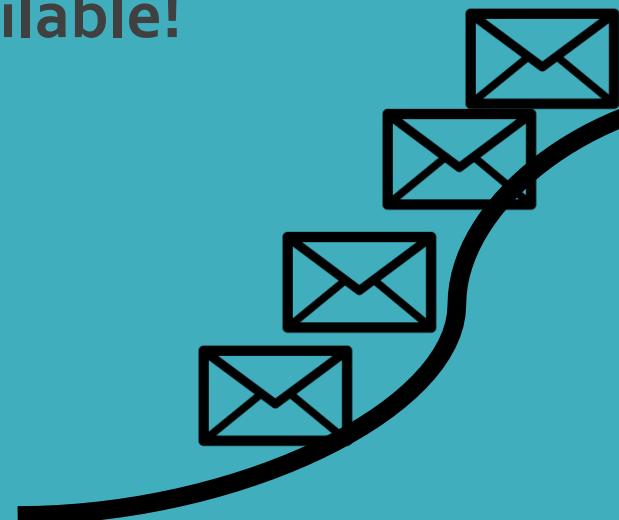
1 operation
per sec



Push Model

Increase buffer size... while you have memory available!

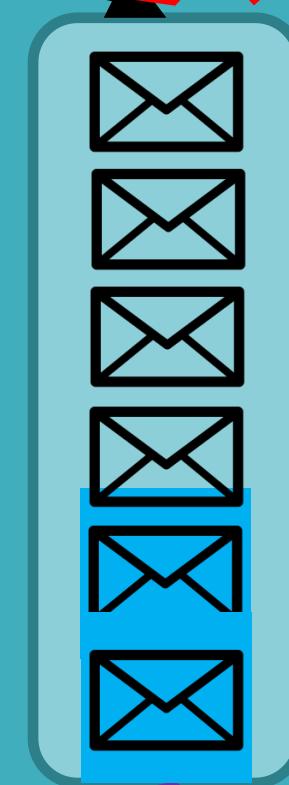
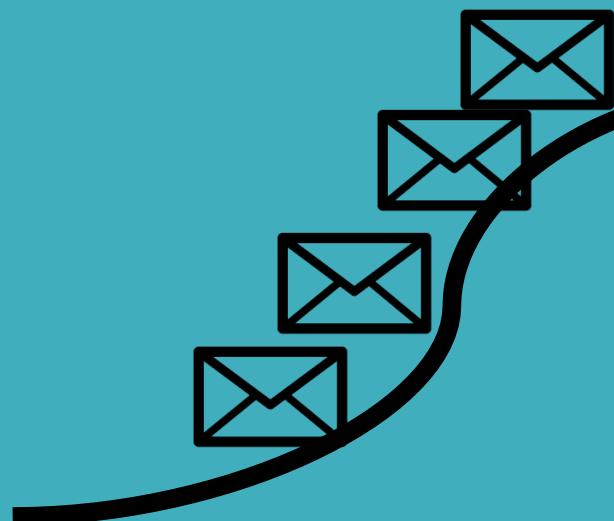
5 operations
per sec



Push Model

Out-Of-Memory

5 operations
per sec



memory increase

1 operation
per sec



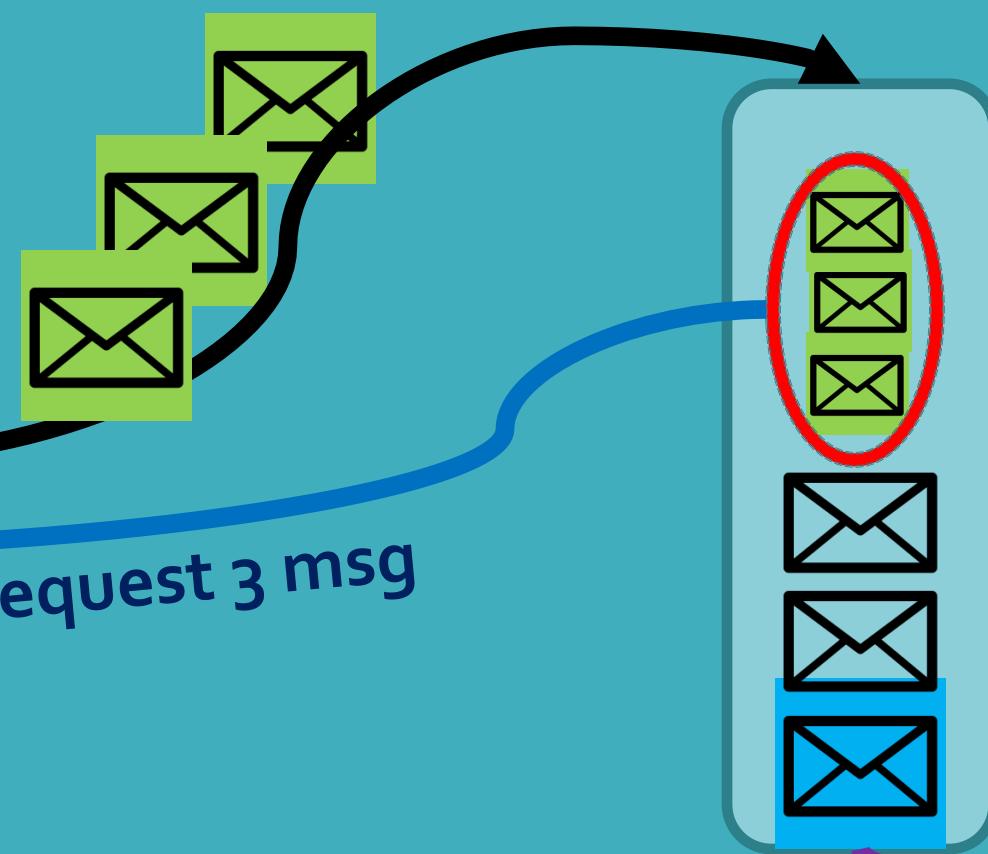
Dynamic Push-Pull

Pull based backpressure

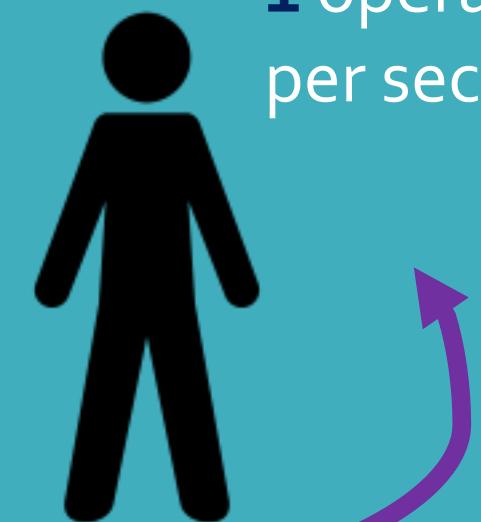
5 operations
per sec

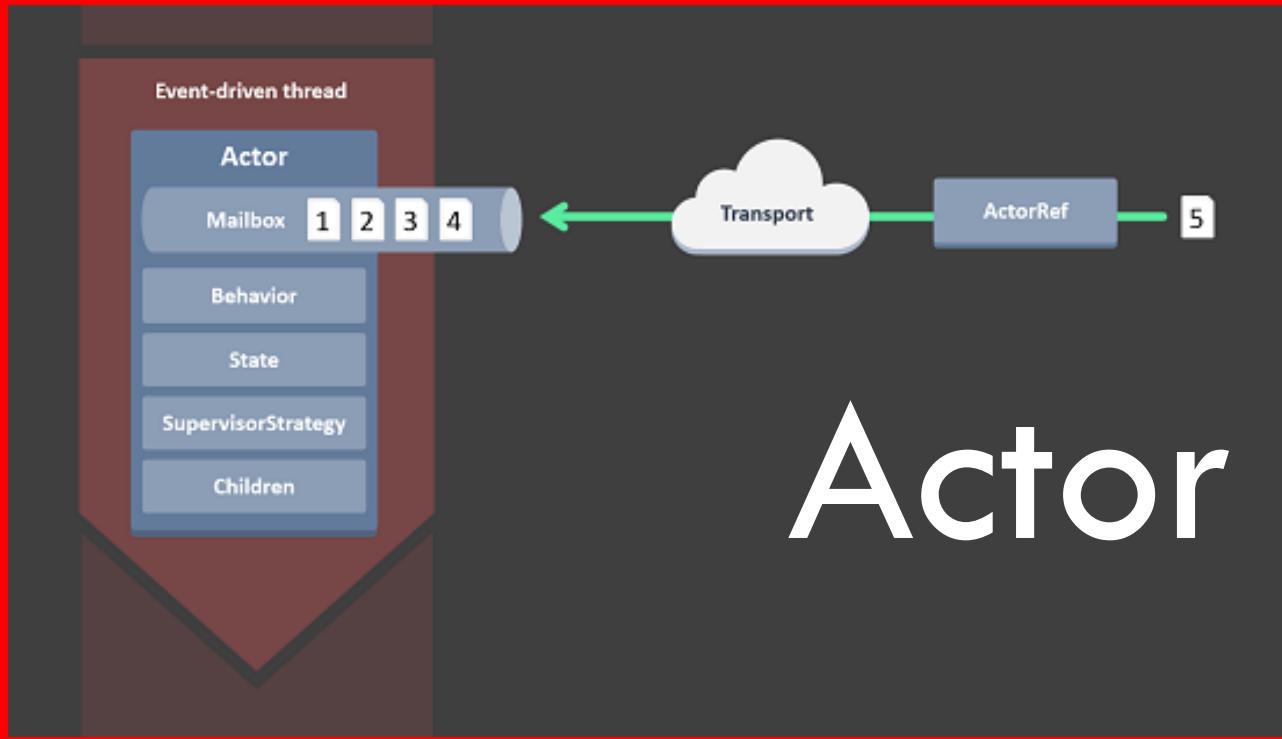


Request 3 msg



1 operation
per sec





Actor Model

Actor is a single and independent unit of computation



Actor Model Message-Driven



What is an Actor?



- **Share Nothing**
- **Message are passed by value**
- Light weight processes/threads communicating through messaging
- Communication only by messages
- Lightweight object
- Processing
- Storage – State
- Running on it's own thread.
- Messages are buffered in a “mailbox”

Actor



KEEP
CALM
AND
LET IT
CRASH

What an Actor Offers

High
throughput

Scale Out & Up

Fault tolerance

High
Availability

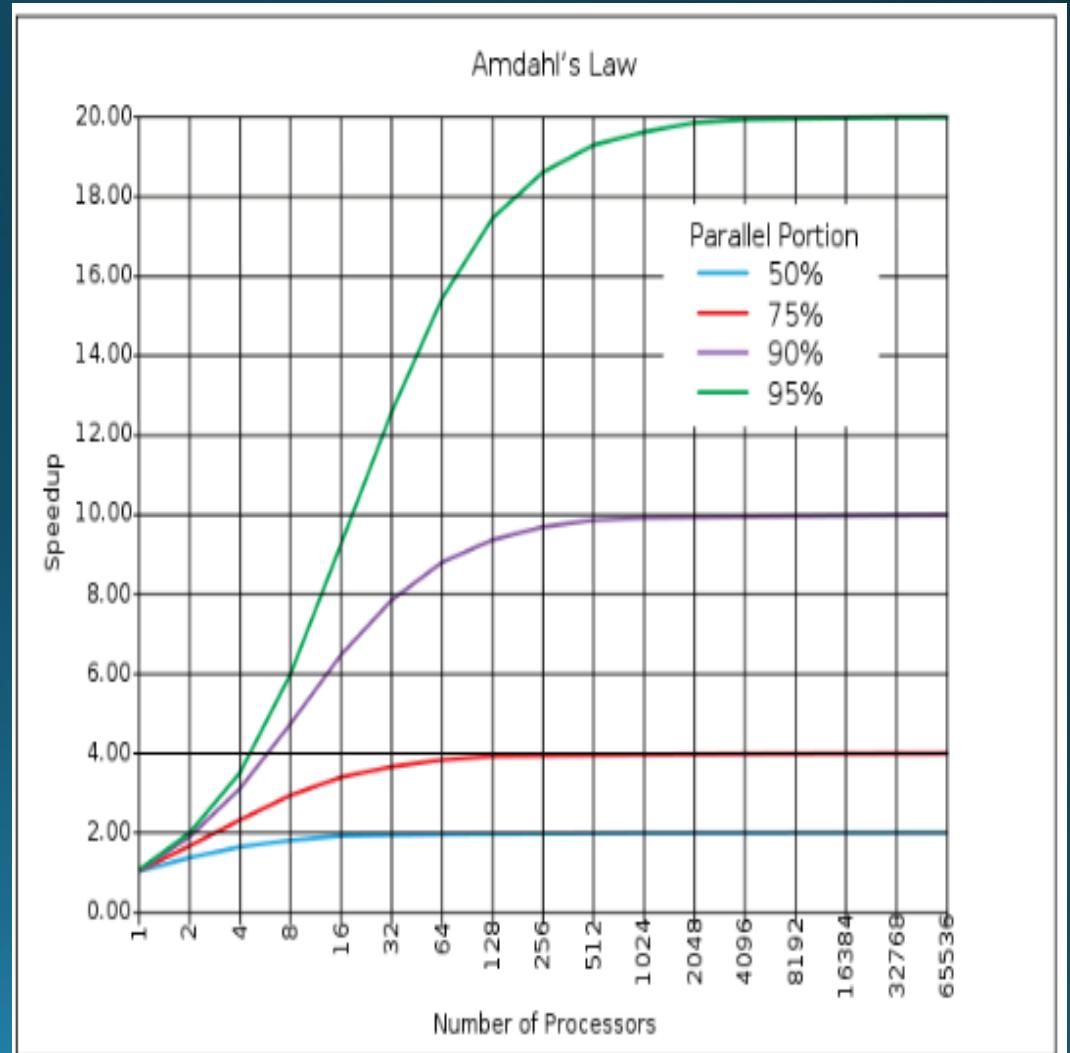
No manual
thread
management

Asynchronous

The new reality : Amdahl's law

The speed-up of a program using multiple processors in parallel computing is limited by the sequential fraction of the program.

For example if 95% of the program can be parallelized, the theoretical maximum speedup using parallel computing would be 20 times, no matter how many processors are used



From Actors to Reactive Streams

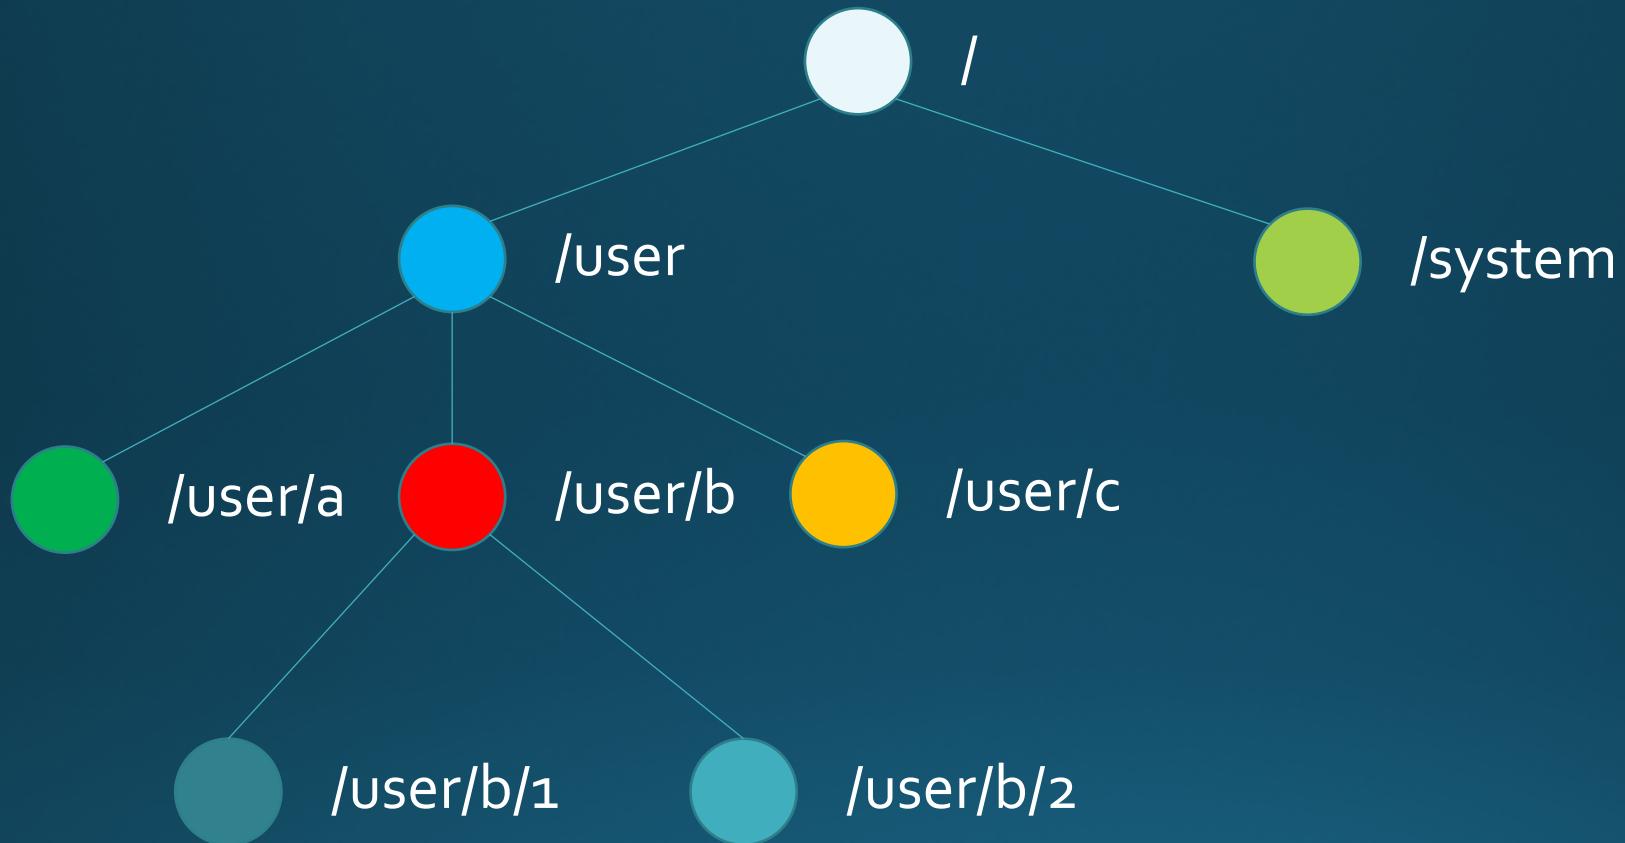
From Actors to Streams

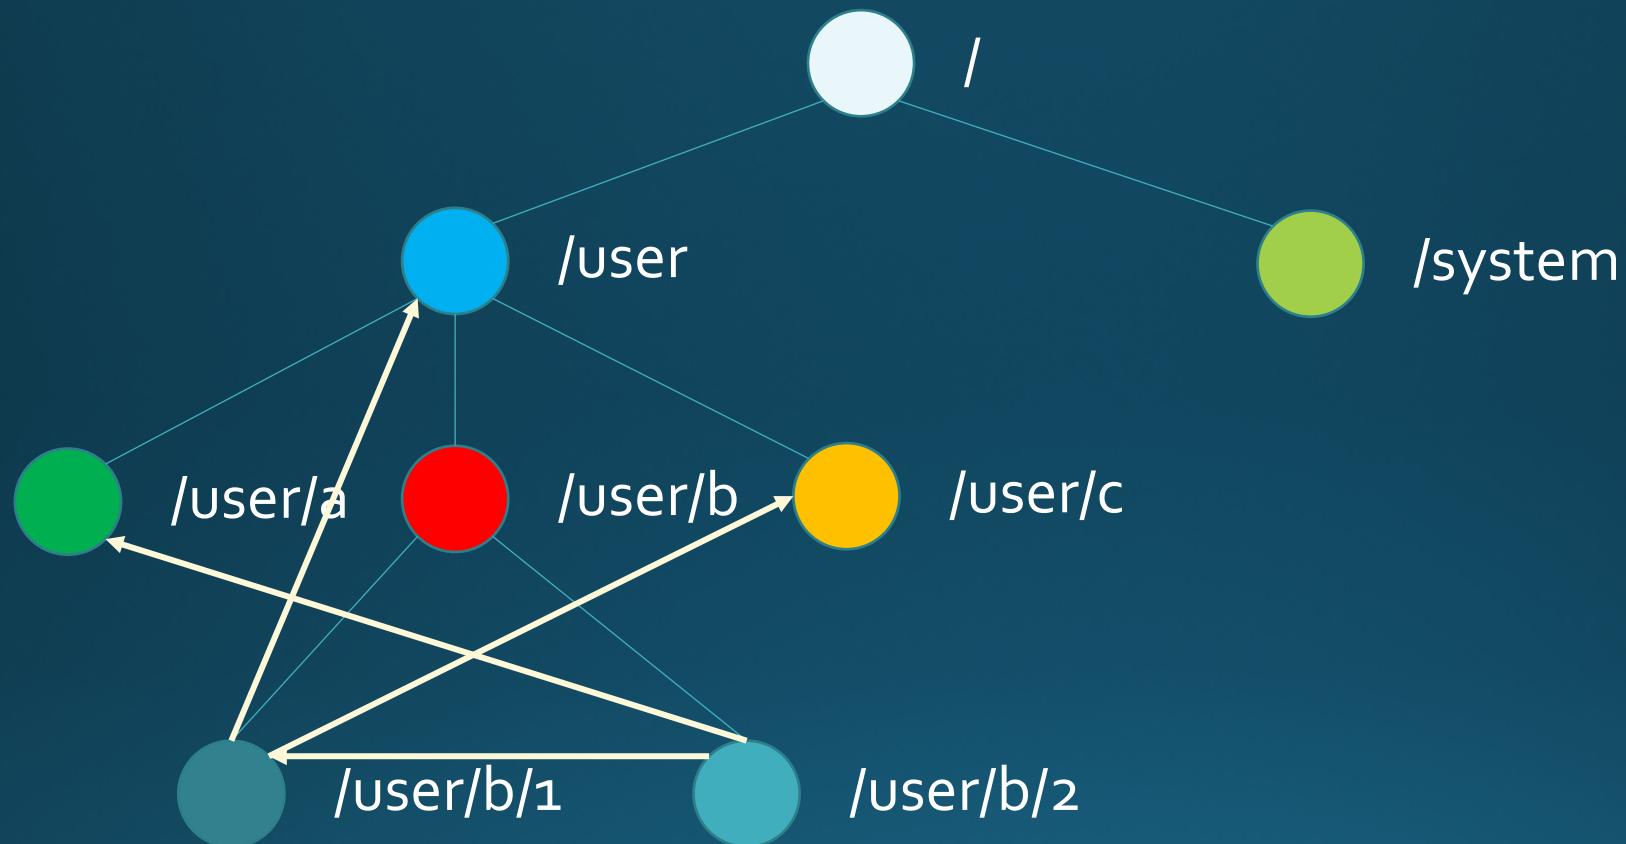
Claim: “actors do not compose”

“By default actors hard-code the receiver of any messages they send. If I create an actor A that sends a message to actor B, and you want to change the receiver to actor C you are basically out of luck.”

Noel Welsh (Underscore.io)

<http://tinyurl.com/AkkaStreamsNdc2>





Claim: “actors do not compose”

“Taken at face value (looking at the API) actors are functions which do not return anything: in Akka both tell() and Receive return Unit. This is why the statement arises that they cannot be composed, and what is meant is that they do not compose like normal functions.”

Roland Kuhn (Akka Tech Lead)

<http://tinyurl.com/AkkaStreamsNdc3>

Actor implementation class in C#

```
public class SampleActor : ReceiveActor
{
    public SampleActor()
    {
        Idle();
    }

    protected override void PreStart() { /* ... */ }

    private void Idle()
    {
        Receive<Job>(job => /* ... */);
    }

    private void Working()
    {
        Receive<Cancel>(job => /* ... */);
    }
}
```


Like in LINQ...

```
var results = db.Companies
    .Join(db.People,
        c => c.CompanyID,
        p => p.PersonID,
        (c, p) => new { c, p })
    .Where(z => z.c.Created >= fromDate)
    .OrderByDescending(z => z.c.Created)
    .Select(z => z.p)
    .ToList();
```

... or pipeline of functions

```
HttpGet pageUrl  
|> fun s -> Regex.Replace(s, "[^A-Za-z']", " ")  
|> fun s -> Regex.Split(s, " ")  
|> Set.ofArray  
|> Set.filter (fun word -> not (Spellcheck word))  
|> Set.iter (fun word -> printfn "%s" word)
```

How about that?

```
val in = Source(1 to 10)
val out = Sink.ignore
val bcast = builder.add(Broadcast[Int](2))
val merge = builder.add(Merge[Int](2))
val f1,f2,f3,f4 = Flow[Int].map(_ + 10)
```

```
source ~> f1 ~> bcast ~> f2 ~> merge ~> f3 ~> sink
          bcast ~> f4 ~> merge ~>
```

Source: <http://doc.akka.io/docs/akka/2.4/scala/stream/stream-graphs.html>

Reactive Streams

Let's unlearn how to create individual actors

...and learn instead high level composition primitives

...that will internally create and connect instances of actors

Example in C#

```
var runnable =  
    Source  
        .From(Enumerable.Range(1, 1000))  
        .Via(Flow.Create<int>().Select(x => x * 2)  
        .To(Sink.ForEach<int>(x => Console.WriteLine(x.ToString)));
```

Graph DSL in C#

```
var graph = GraphDsl.Create(builder =>
{
    var bcast = builder.Add(new Broadcast<int>(2));
    var merge = builder.Add(new Merge<int, int>(2));
    var count = Flow.FromFunction(new Func<int, int>(x => 1));
    var sum = Flow.Create<int>().Sum((x, y) => x + y);
    builder.From(bcast.Out(0)).To(merge.In(0));
    builder.From(bcast.Out(1)).Via(count).Via(sum).To(merge.In(1));
    return new FlowShape<int, int>(bcast.In, merge.Out);
});
```

Reactive Streams

“Reactive Streams is an initiative to provide a standard for asynchronous stream processing with non-blocking back pressure. This encompasses efforts aimed at runtime environments (JVM and JavaScript) as well as network protocols

<http://www.reactive-streams.org>

Participants in Reactive Stream initiative

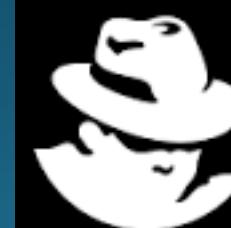
NETFLIX

ORACLE®

Lightbend

Pivotal

twitter



redhat.

Reactive Streams

“Reactive Streams is an initiative to provide a standard for **asynchronous** stream processing with non-blocking back pressure. This encompasses efforts aimed at runtime environments (JVM and JavaScript) as well as network protocols

asynchronous

<http://www.reactive-streams.org>



Is all about Asynchronicity !

Reactive Streams

“Reactive Streams is an initiative to provide a standard for **asynchronous stream processing** with non-blocking back pressure. This encompasses efforts aimed at runtime environments (JVM and JavaScript) as well as network protocols

asynchronous
stream processing

<http://www.reactive-streams.org>

Reactive Streams

“Reactive Streams is an initiative to provide a standard for **asynchronous stream processing** with **non-blocking back pressure**. This encompasses efforts aimed at runtime environments (JVM and JavaScript) as well as network protocols

<http://www.reactive-streams.org>

asynchronous
stream processing
non-blocking back pressure

What do Reactive Streams solve?



Getting Data across an
Asynchronous Boundary
without running **Out-Of-Memory**

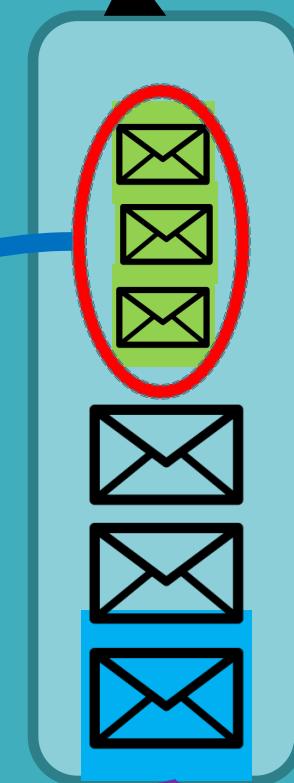
Dynamic Push-Pull

Pull based backpressure

5 operations
per sec



Request 3 msg

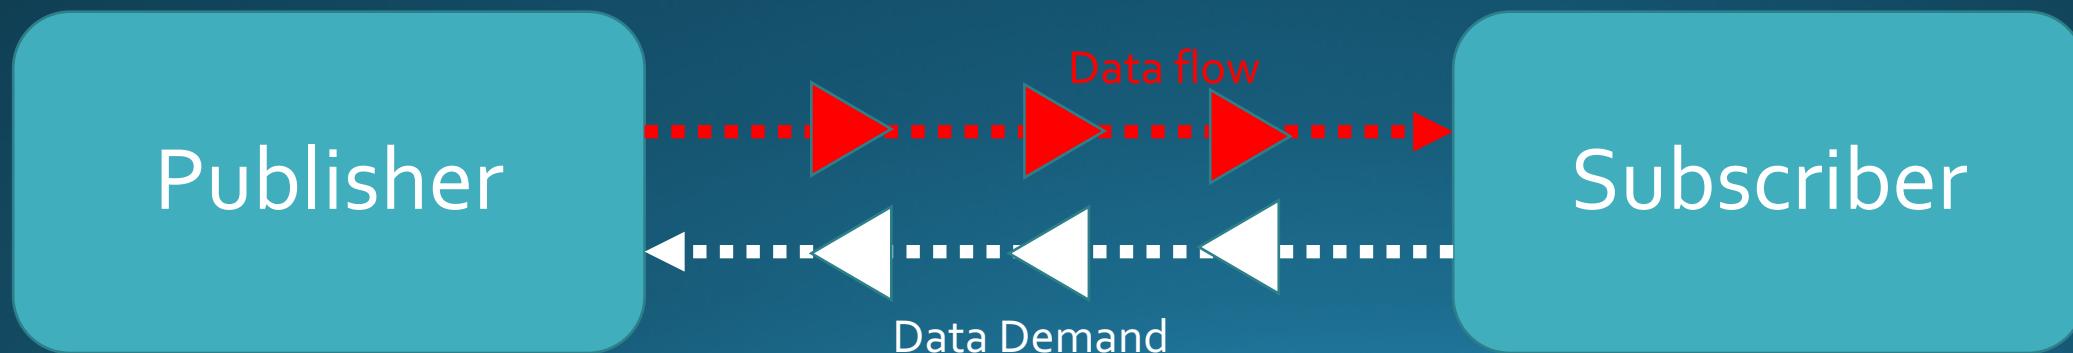


1 operation
per sec

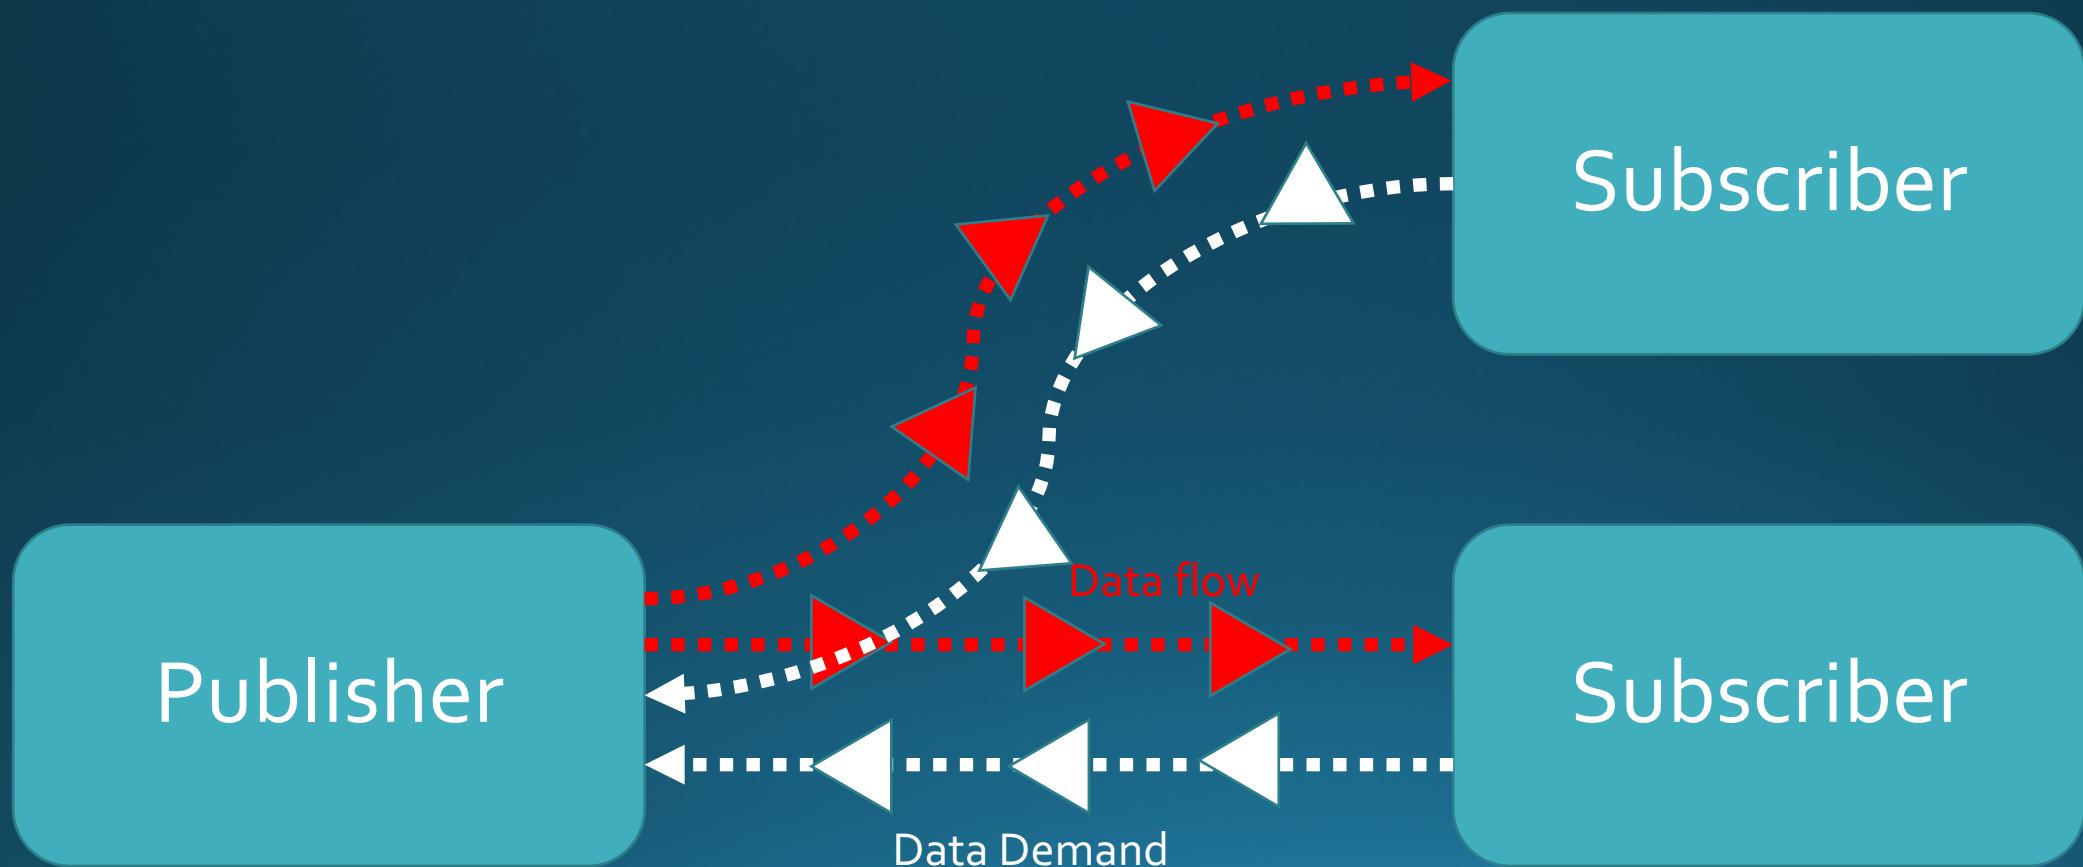


Push and Pull

- data items flow downstream
- data demand flows upstream
- data items flow only when there is demand
 - recipient is in control of incoming data rate
 - data in flight is bounded by signaled demand
 - “push” behavior when consumer is faster
 - “pull” behavior when producer is faster



Splitting and merging the data flow



You might consider Akka Streams if...

- You have used Akka but want to save yourself from details of actor definition and management
- You are familiar with Reactive Streams and looking for ready-for-use implementations of its specification
- You find Akka Streams stages building blocks suitable to compose your process workflows
- You want to take advantage of Akka Streams backpressure to have a fine grained control over throughput of your process workflow stages

Reactive Stream interfaces

```
public interface IPublisher<out T>
{
    void Subscribe(ISubscriber<T> subscriber);
}

public interface ISubscriber<in T>
{
    void OnSubscribe(Subscription subscription);
    void OnNext(T element);
    void OnError(Exception cause);
    void OnComplete();
}

public interface Subscription
{
    void Request(long n);
    void Cancel();
}
```

Connection

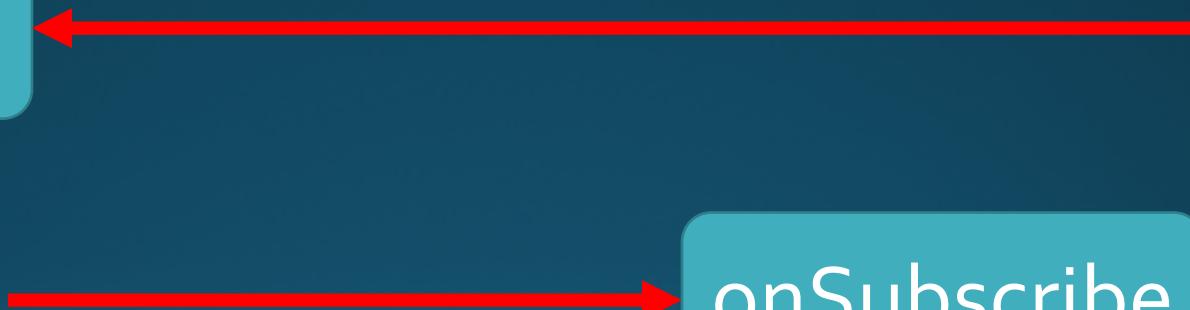
▪ Publisher

Subscribe

Subscriber

Subscription

onSubscribe



Flow

▪ Publisher

request

Elements

request

request

Elements

▪ Subscriber

onNext

onNext



Complete

▪ Publisher



▪ Subscriber



Elements



Failure

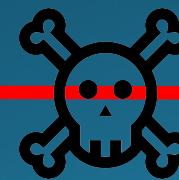
: Publisher



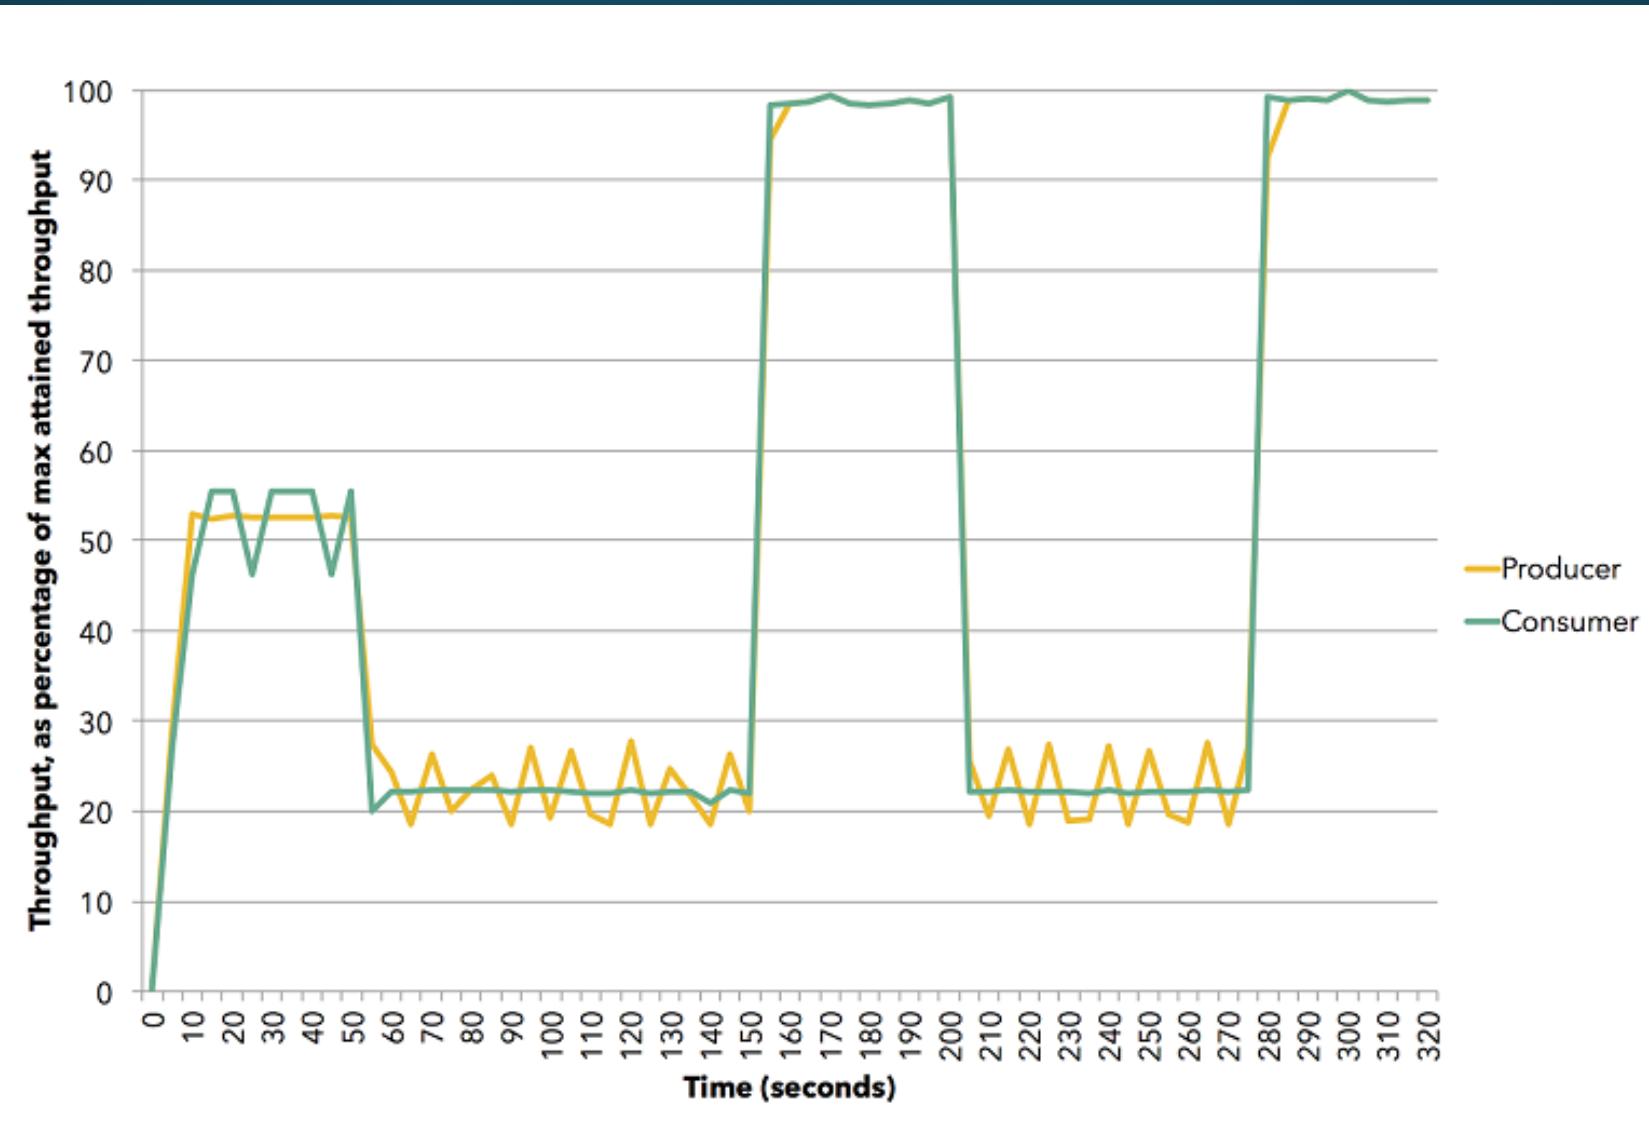
: Subscriber



Elements



onError



Source: <https://data-artisans.com/blog/how-flink-handles-backpressure> (Apache Flink)

Akka Streams

- *Streams complement Actors, they do not replace them.*
- *Actors – distribution (location transparency)*
- *Streams – back-pressured*
- Streams talking to **Actors** & **Actors** talking to **Streams**



Push-Pull bidirectional communication

Akka Stream Lexicon

Source

This is where the pipeline starts. A Source takes data from input and has a single output to be written into.

Sink

The pipeline ends here. A Sink has a single input to be written into.

Flow

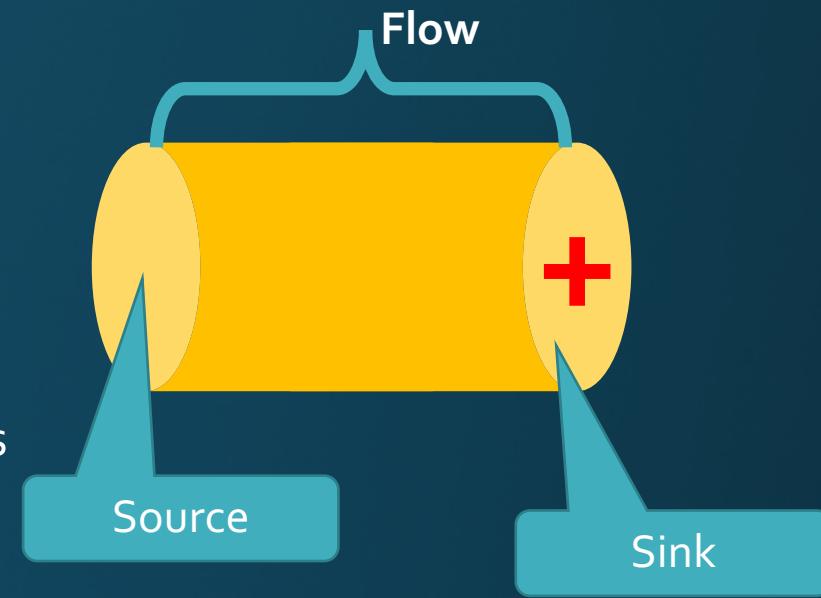
This is the basic abstraction where transformation of data takes place. A Flow has one input and one output

RunnableGraph

The entire topology of (Source -> Flow* -> Sink) a graph in Akka Streams.

Materializer

A materializer defines how the transformations are converted into asynchronous processes.



Source stages in Akka.NET Streams

- FromEnumerable
- From
- Single
- Repeat
- Cycle
- Tick
- FromTask
- Unfold
- UnfoldAsync
- Empty
- Maybe
- Failed
- ActorPublisher
- ActorRef
- Combine
- UnfoldResource
- Queue
- AsSubscriber
- FromPublisher
- ZipN
- ZipWithN
- FromInputStream
- AsOutputStream
- FromFile

Sink stages in Akka.NET Streams

- First
- FirstOrDefault
- Last
- LastOrDefault
- Ignore
- Cancelled
- Seq
- Foreach
- ForeachParallel
- OnComplete
- Aggregate
- Sum
- Combine
- ActorRef
- ActorRefWithAck
- ActorSubscriber
- AsPublisher
- FromSubscriber
- FromOutputStream
- AsInputStream
- ToFile

Flow stages in Akka.NET Streams

- Select
- SelectMany
- StatefulSelectMany
- Where
- Collect
- Grouped
- Sliding
- Scan
- Aggregate
- Skip
- SkipWhile
- Take
- TakeWhile
- Recover
- RecoverWith
- Detach
- Throttle
- SelectAsync
- SelectAsyncUnordered
- TakeWithin
- SkipWithin
- GroupedWithin
- InitialDelay
- Delay
- Conflate
- ConflateWithSeed
- Batch
- BatchWeighted
- Expand
- Buffer
- PrefixAndTail
- GroupBy
- SplitWhen
- SplitAfter
- ConcatMany
- MergeMany
- InitialTimeout
- CompletionTimeout
- IdleTimeout
- BackpressureTimeout
- KeepAlive
- InitialDelay
- Merge
- MergeSorted
- Zip
- ZipWith
- ZipWithIndex
- Concat
- Prepend
- OrElse
- Interleave
- Unzip
- UnzipWith
- Broadcast
- Balance
- WatchTermination
- Monitor
- InitialTimeout

Stream materialization

- Stream declaration prepares a blueprint (an execution plan)
- Execution begins with stream materialization
- Only runnable graphs can be materialized
- Multiple stream stages can be allocated to a single actor
- Most of flow stages preserve input order of elements

Compare with LINQ statements

```
var results = db.Companies
    .Join(db.People,
        c => c.CompanyID,
        p => p.PersonID,
        (c, p) => new { c, p })
    .Where(z => z.c.Created >= fromDate)
    .OrderByDescending(z => z.c.Created)
    .Select(z => z.p)
    .ToList();
```

Materializing a stream

```
var runnable =  
    Source  
        .From(Enumerable.Range(1, 1000))  
        .Via(Flow.Create<int>().Select(x => x * 2))  
        .To(Sink.ForEach<int>(x => Console.WriteLine(x.ToString)));
```

Materializing a stream

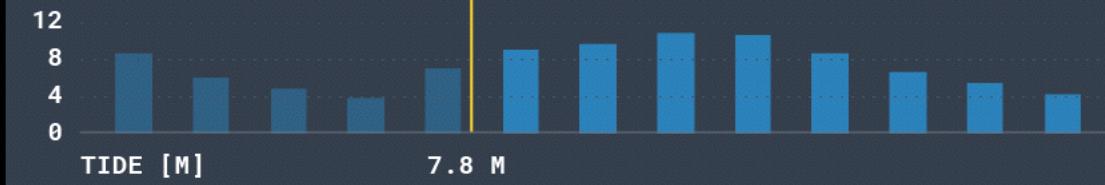
```
var runnable =  
    Source  
        .From(Enumerable.Range(1, 1000))  
        .Via(Flow.Create<int>().Select(x => x * 2))  
        .To(Sink.ForEach<int>(x => Console.WriteLine(x.ToString)));  
  
var system = ActorSystem.Create("MyActorSystem");  
using (var materializer = ActorMaterializer.Create(system))  
{  
    await runnable.Run(materializer);  
}
```


TOTAL**286 MW**

Active Power

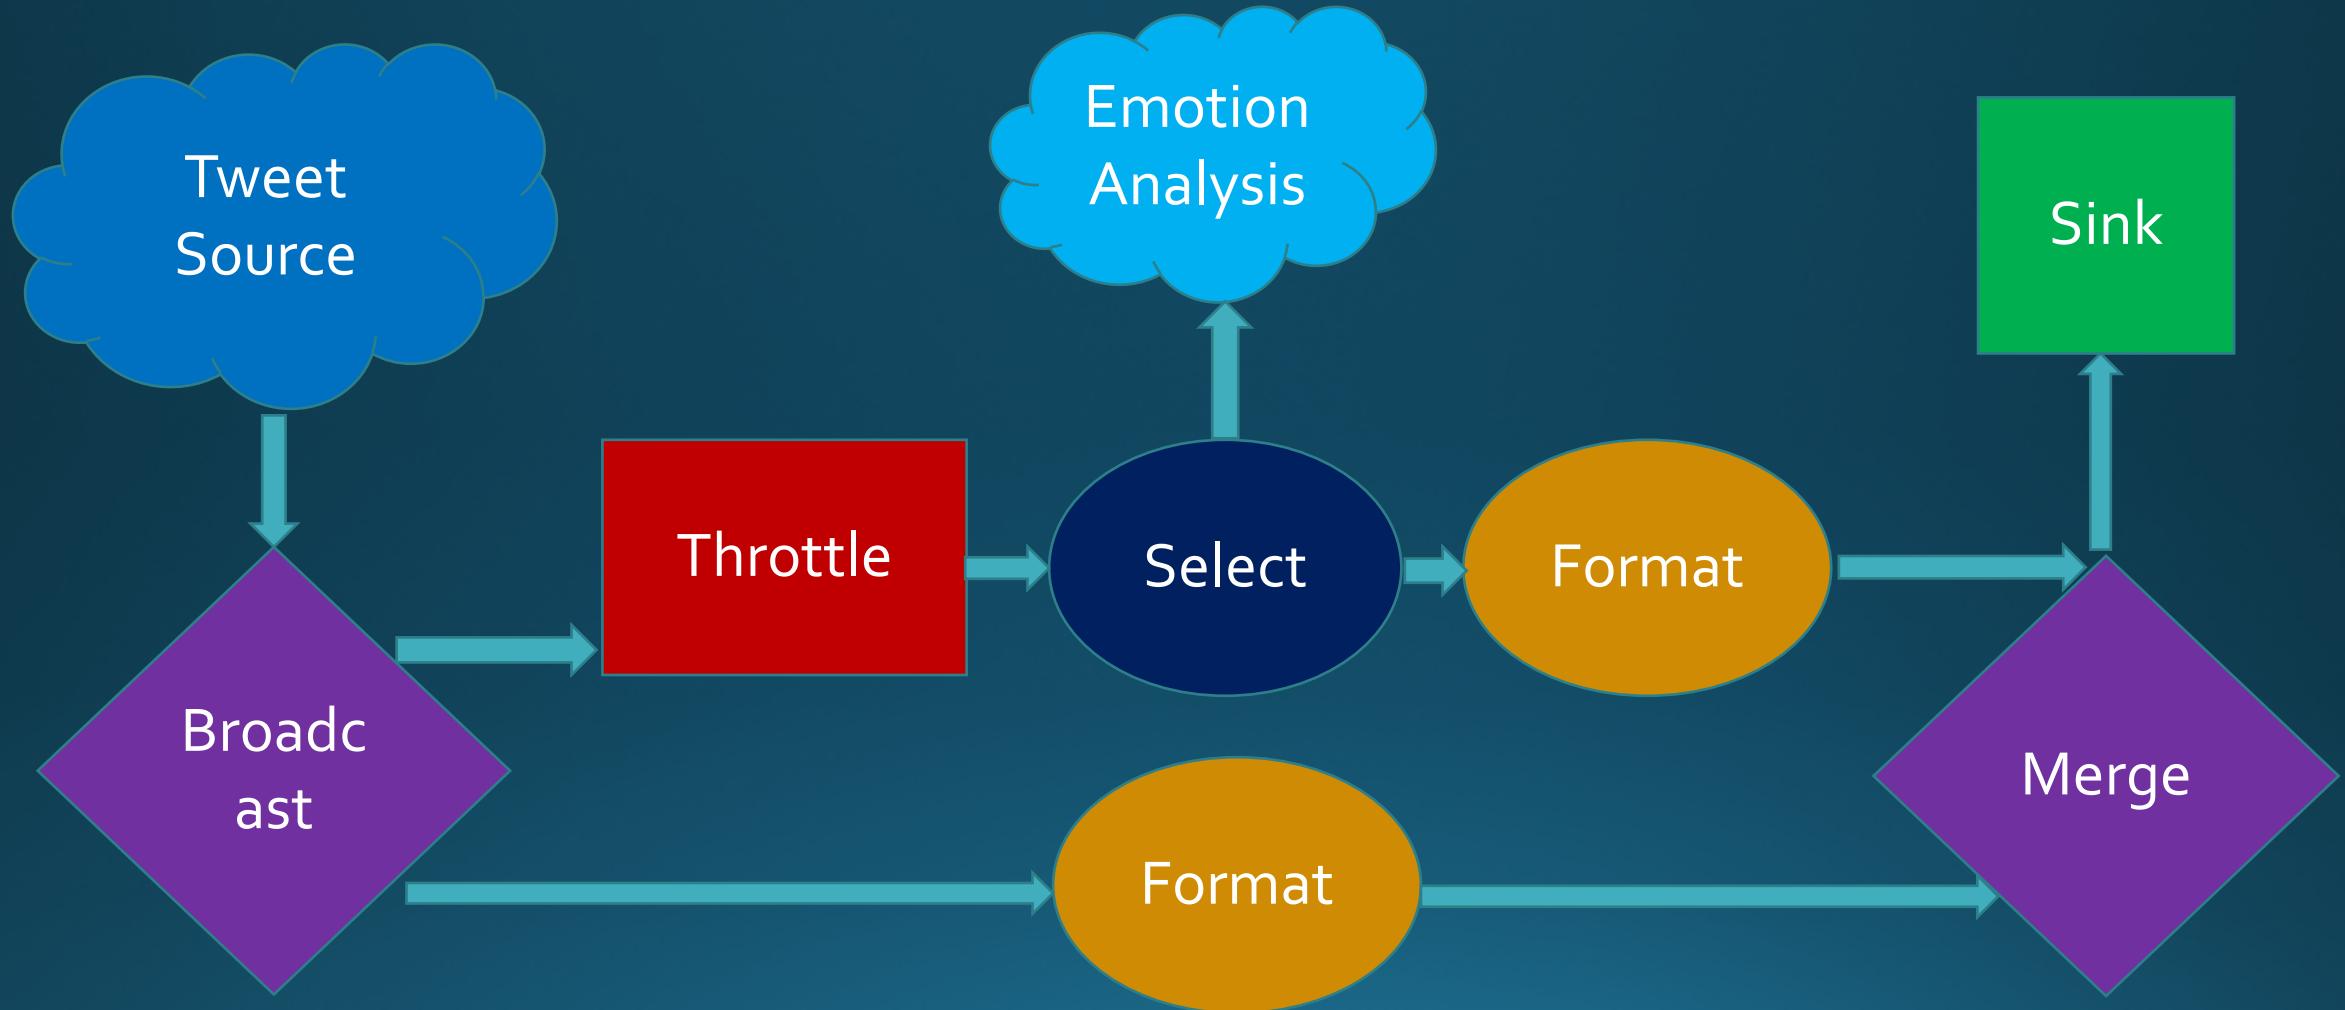
**9.1 °C**
Air Temp**14.6 °C**
Water Temp

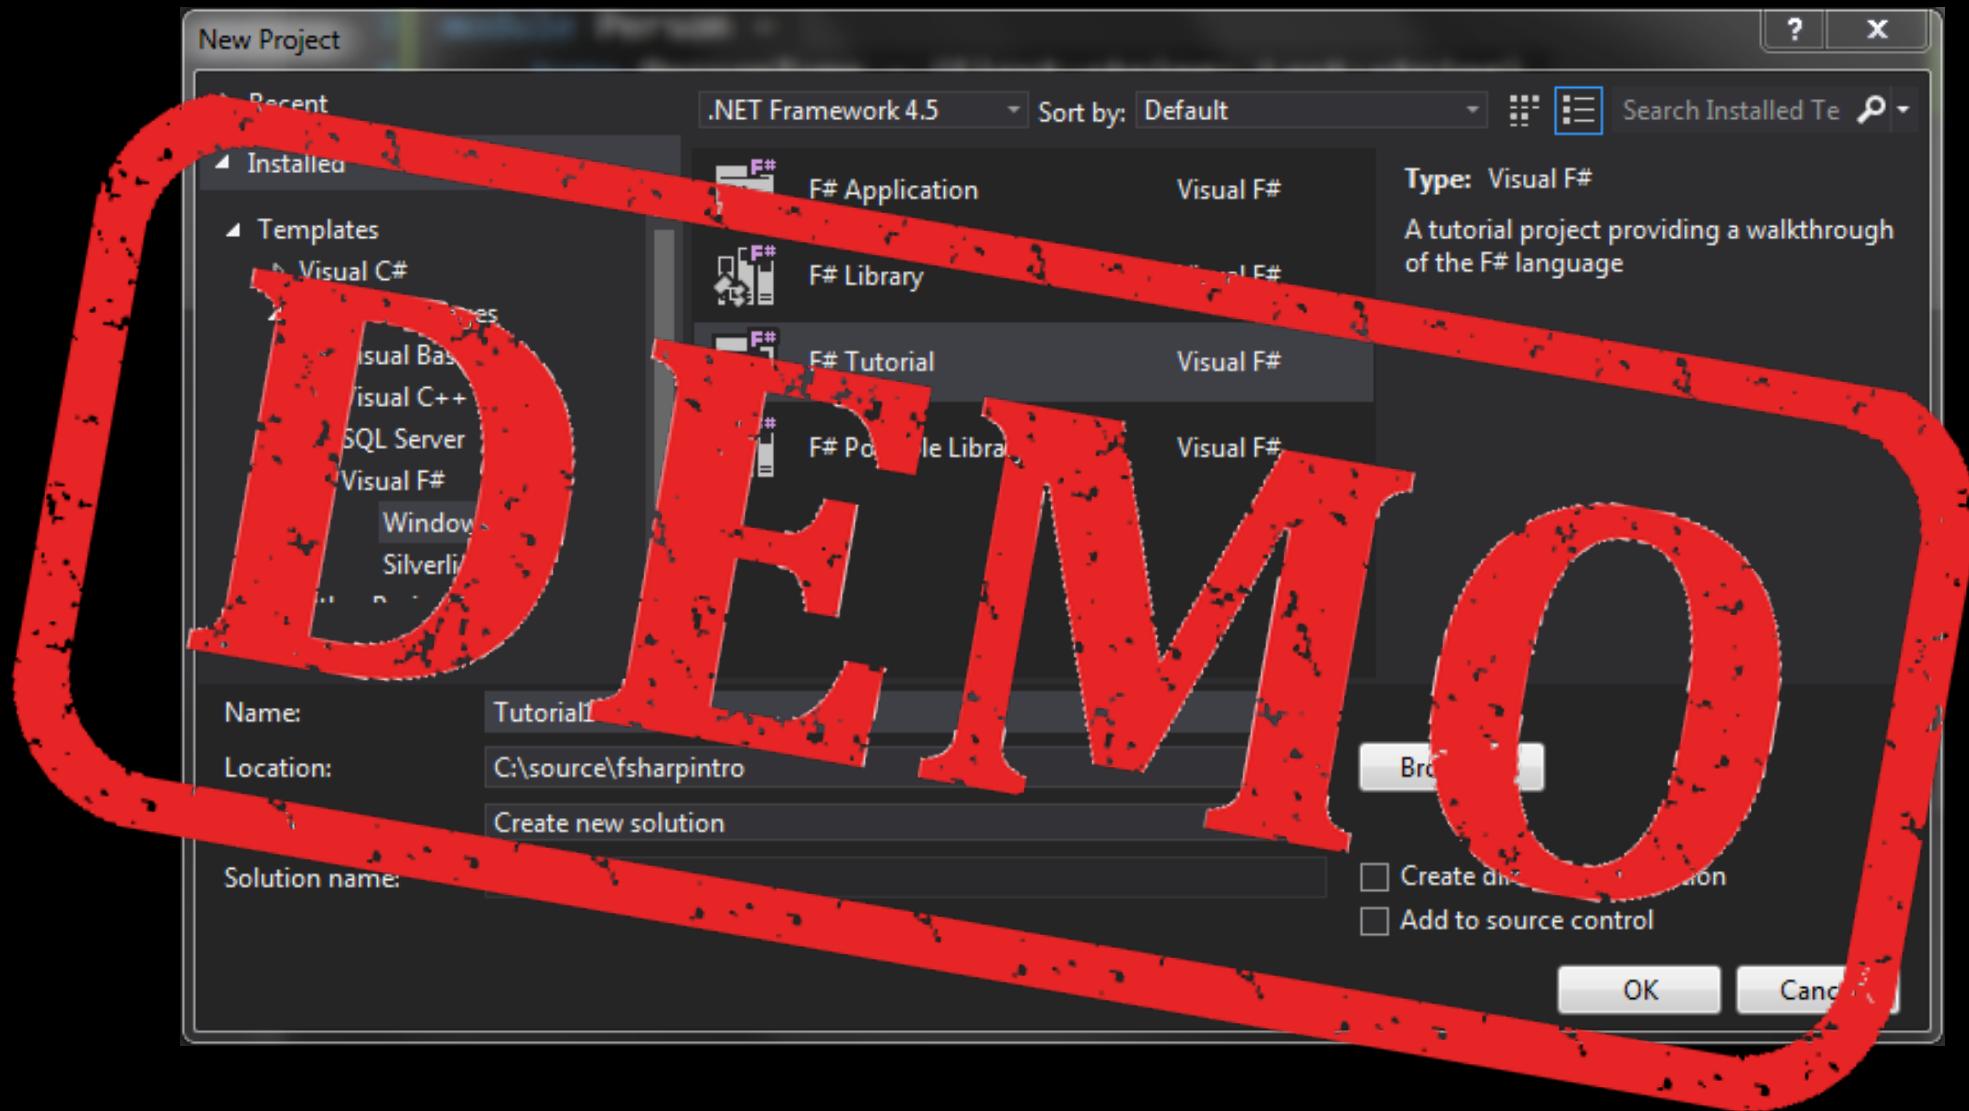
FR 03 06 09 12 15 18 21 SA 03 06 09 12



Reactive Tweets

- Both Akka (Scala/JVM) and Akka.NET contain samples showing a stream of Twitter tweets
- Akka.NET sample is based on Tweetinvi (<https://github.com/linvil/tweetinvi>)
- Tweetinvi expose several streams of Twitter data that don't support Reactive Streams
 - IFilteredStream
 - ISampleStream
 - ITrackedStream
 - ITweetStream
 - IUserStream
- Reactive Tweets use Source.ActorRef built-in primitive to emit elements from Tweetinvi streams to Akka Stream source





That's all Folks!